

## CS 4348/5348 Operating Systems, Homework #2

1. Implement a monitor solution for the bakery problem. Assume that you have N salesmen and customers arrive at arbitrary times. Note: neither the customers nor the salesmen should have a busy waiting.

Monitor

```
{ int sales_count, initialized to N
  int sales_wait, initialized to 0
  int cust_count, initialized to 0
  condition sales, cust;

  get_service()
  { if (sales_count <= 0) cust.wait;
    cust_count = cust_count + 1;
    if (sales_count = sales_wait) sales.signal;
    sales_count = sales_count - 1;
  }

  release_service()
  { } -- this can be omitted

  prepare_service ()
  { if (cust_count <= 0) { sales_wait = sales_wait + 1; sales.wait; sales_wait = sales_wait - 1; }
    cust_count = cust_count - 1;
  }

  complete_service()
  { sales_count = sales_count + 1;
    cust.signal;
  }
} // end monitor

customer
{ get_service ();
  receive the service
  release_service ();
}

salesman
{ prepare_service();
  provide the service
  complete_service();
}
```

2. Discuss the steps for implementing the Monitor construct using integer semaphores. Consider letting the signalee continue after condition signaling. You need to consider all steps in the implementation and give more implementation details than those shown in lecture slides (which are incomplete).

The monitor should be protected by a semaphore “mutex”.

For each condition cond, we need to use a semaphore and a separate counter (let condsem and counter represent them). For cond.wait, we always set condsem.count to 0 (e.g., initialize semaphore again to force it to be set to 0) before wait(condsem), so that cond.wait will be forced to wait. For cond.signal, we should set condsem.count to negative of the counter value before signal(condsem) (so that a process in condsem.queue can be released correctly or when condsem.queue is empty, condsem.count = 0).

We also need to use a semaphore temp and maintain a counter tempcount for holding the signalers after conditional signals. When implementing cond.signal, after signal(condsem), execute  
tempcount++; wait(temp);

When any process is exiting Monitor, execute:

if tempcount > 0 { tempcount--; temp.signal; } else mutex.signal;

- Given the system state, follow Banker's algorithm to determine whether to grant or deny the requests. You need to compute the "Needed" matrix (numbers of resources the processes still need) and the "Available" vector (the available resources currently in the system) first. Then you need to pretend to allocate and perform the safety analysis. You need to show your work for all the steps.

|       | Total | Allocated |       |       |       | MaxReq |       |       |       | Needed |       |       |       | Available |
|-------|-------|-----------|-------|-------|-------|--------|-------|-------|-------|--------|-------|-------|-------|-----------|
|       |       | $P_1$     | $P_2$ | $P_3$ | $P_4$ | $P_1$  | $P_2$ | $P_3$ | $P_4$ | $P_1$  | $P_2$ | $P_3$ | $P_4$ |           |
| $R_1$ | 3     | 1         | 0     | 0     | 0     | 1      | 1     | 2     | 2     |        |       |       |       |           |
| $R_2$ | 6     | 1         | 1     | 0     | 1     | 2      | 2     | 2     | 1     |        |       |       |       |           |
| $R_3$ | 4     | 0         | 0     | 1     | 1     | 2      | 1     | 1     | 2     |        |       |       |       |           |

|       | Requested |       |       |       |
|-------|-----------|-------|-------|-------|
|       | $P_1$     | $P_2$ | $P_3$ | $P_4$ |
| $R_1$ | 0         | 1     | 0     | 1     |
| $R_2$ | 0         | 0     | 0     | 0     |
| $R_3$ | 1         | 1     | 0     | 0     |

|       | Total | Allocated |       |       |       | MaxReq |       |       |       | Needed |       |       |       | Available |
|-------|-------|-----------|-------|-------|-------|--------|-------|-------|-------|--------|-------|-------|-------|-----------|
|       |       | $P_1$     | $P_2$ | $P_3$ | $P_4$ | $P_1$  | $P_2$ | $P_3$ | $P_4$ | $P_1$  | $P_2$ | $P_3$ | $P_4$ |           |
| $R_1$ | 3     | 1         | 0     | 0     | 0     | 1      | 1     | 2     | 2     | 0      | 1     | 2     | 2     | 2         |
| $R_2$ | 6     | 1         | 1     | 0     | 1     | 2      | 2     | 2     | 1     | 1      | 1     | 2     | 0     | 3         |
| $R_3$ | 4     | 0         | 0     | 1     | 1     | 2      | 1     | 1     | 2     | 2      | 1     | 0     | 1     | 2         |

#Preliminary Checks

Request is Request[i, j] < Need[i, j] for any i, j  $\Rightarrow$  no problem

Request is  $\sum_i \text{Request}[i, j] < \text{Available}[j]$  for any j  $\Rightarrow$  no problem

|       | Total | Allocated |       |       |       | MaxReq |       |       |       | Needed |       |       |       | Available |
|-------|-------|-----------|-------|-------|-------|--------|-------|-------|-------|--------|-------|-------|-------|-----------|
|       |       | $P_1$     | $P_2$ | $P_3$ | $P_4$ | $P_1$  | $P_2$ | $P_3$ | $P_4$ | $P_1$  | $P_2$ | $P_3$ | $P_4$ |           |
| $R_1$ | 3     | 1         | 1     | 0     | 1     | 1      | 1     | 2     | 2     | 0      | 0     | 2     | 1     | 0         |
| $R_2$ | 6     | 1         | 1     | 0     | 1     | 2      | 2     | 2     | 1     | 1      | 1     | 2     | 0     | 3         |
| $R_3$ | 4     | 1         | 1     | 1     | 1     | 2      | 1     | 1     | 2     | 1      | 0     | 0     | 1     | 0         |

$P_2$  can continue, choose  $P_2$  to continue and  $P_2$  returns all its resources

|       | Total | Allocated |       |       |       | MaxReq |       |       |       | Needed |       |       |       | Available |
|-------|-------|-----------|-------|-------|-------|--------|-------|-------|-------|--------|-------|-------|-------|-----------|
|       |       | $P_1$     | $P_2$ | $P_3$ | $P_4$ | $P_1$  | $P_2$ | $P_3$ | $P_4$ | $P_1$  | $P_2$ | $P_3$ | $P_4$ |           |
| $R_1$ | 3     | 1         | 0     | 0     | 1     | 1      | 1     | 2     | 2     | 0      | 0     | 2     | 1     | 1         |
| $R_2$ | 6     | 1         | 0     | 0     | 1     | 2      | 2     | 2     | 1     | 1      | 0     | 2     | 0     | 4         |
| $R_3$ | 4     | 1         | 0     | 1     | 1     | 2      | 1     | 1     | 2     | 1      | 0     | 0     | 1     | 1         |

$P_4$  or  $P_1$  can continue, choose  $P_4$  to continue and  $P_4$  returns all its resources

|       | Total | Allocated |       |       |       | MaxReq |       |       |       | Needed |       |       |       | Available |
|-------|-------|-----------|-------|-------|-------|--------|-------|-------|-------|--------|-------|-------|-------|-----------|
|       |       | $P_1$     | $P_2$ | $P_3$ | $P_4$ | $P_1$  | $P_2$ | $P_3$ | $P_4$ | $P_1$  | $P_2$ | $P_3$ | $P_4$ |           |
| $R_1$ | 3     | 1         | 0     | 0     | 0     | 1      | 1     | 2     | 2     | 0      | 0     | 2     | 0     | 2         |
| $R_2$ | 6     | 1         | 0     | 0     | 0     | 2      | 2     | 2     | 1     | 1      | 0     | 2     | 0     | 5         |

$R_3$  4 1 0 1 0 2 1 1 2 1 0 0 0 2  
 $P_1$  or  $P_3$  can continue, choose  $P_1$  to continue and  $P_1$  returns all its resources  
 $R_1$  3 0 0 0 0 1 1 2 2 0 0 2 0 3  
 $R_2$  6 0 0 0 0 2 2 2 1 0 0 2 0 6  
 $R_3$  4 0 0 1 0 2 1 1 2 0 0 0 0 3  
 $P_3$  can continue and all processes can finish  $\Rightarrow$  So, it is safe, grant the request

4. Given the current system state, follow the deadlock detection algorithm to check whether there is a deadlock. You need to show your work clearly. If there is deadlock, you should also indicate which processes are involved in the deadlock.

|       | Total | Allocated |       |       | Requested |       |       |
|-------|-------|-----------|-------|-------|-----------|-------|-------|
|       |       | $P_1$     | $P_2$ | $P_3$ | $P_1$     | $P_2$ | $P_3$ |
| $R_1$ | 1     | 1         | 0     | 0     | 0         | 0     | 1     |
| $R_2$ | 1     | 0         | 1     | 0     | 1         | 0     | 0     |
| $R_3$ | 1     | 0         | 0     | 0     | 0         | 1     | 1     |
| $R_4$ | 1     | 0         | 0     | 1     | 1         | 0     | 0     |

  

|       | Total | Allocated |       |       | Avl | Requested |       |       |                                    |
|-------|-------|-----------|-------|-------|-----|-----------|-------|-------|------------------------------------|
|       |       | $P_1$     | $P_2$ | $P_3$ |     | $P_1$     | $P_2$ | $P_3$ |                                    |
| $R_1$ | 1     | 1         | 0     | 0     | 0   | 0         | 0     | 1     |                                    |
| $R_2$ | 1     | 0         | 1     | 0     | 0   | 1         | 0     | 0     |                                    |
| $R_3$ | 1     | 0         | 0     | 0     | 1   | 0         | 1     | 1     | $P_2$ can get the resources        |
| $R_4$ | 1     | 0         | 0     | 1     | 0   | 1         | 0     | 0     | $\Rightarrow P_2$ return resources |
| ----- |       |           |       |       |     |           |       |       |                                    |
| $R_1$ | 1     | 1         | 0     | 0     | 0   | 0         | 0     | 1     |                                    |
| $R_2$ | 1     | 0         | 0     | 0     | 1   | 1         | 0     | 0     | No process can get the resources   |
| $R_3$ | 1     | 0         | 0     | 0     | 1   | 0         | 0     | 1     | $\Rightarrow$ Deadlock             |
| $R_4$ | 1     | 0         | 0     | 1     | 0   | 1         | 0     | 0     | $P_1$ and $P_3$ involved           |

5. Consider a memory system with 1GB space (=1024MB). Currently the memory free list contains ((128MB, 64MB), (320MB, 8MB), (512MB, 16MB), (768MB, 256MB)) where each item (x, y) in the list represents the starting location x and length y of a free partition. Four jobs  $J_1$ ,  $J_2$ ,  $J_3$ , and  $J_4$  of sizes 14MB, 20MB, 160MB, and 64MB, respectively, are to be loaded to the memory. Consider the first fit allocation policy and answer the following questions.

- (a) Consider each job  $J_i$ ,  $1 \leq i \leq 4$ . When  $J_i$  is running, what would be the value for the base register and bound register?

$J_1$  gets (128MB, 64MB) slot

$\Rightarrow$  base register 0x08000000, bound register 0x08dfffff

(or 0x08e00000 if we use  $A \geq B$  as the circuit, but consider the above as the correct answer)

$J_2$  gets allocated after  $J_1$  in (128MB, 64MB) slot

$\Rightarrow$  base register 0x08e00000, bound register 0x0a1fffff (starting address of next slot is 0x0a200000)

$J_3$  gets (768MB, 256MB) slot

$\Rightarrow$  base register 0x30000000, bound register 0x39ffffff (starting address of next slot is 0x3a000000)

$J_4$  gets allocated after  $J_3$  in (768MB, 256MB) slot

$\Rightarrow$  base register 0x3a000000, bound register 0x3dffffff (starting address of next slot is 0x3e000000)

(b) Now  $J_4$  is running and it accesses the logical address 0x04213a0c. Compute the corresponding physical address. Also, check and determine whether this causes access violation.

$J_4$  base register 0x3a000000 + logical address 0x04213a0c  
 $\Rightarrow$  physical address 0x3e213a0c > bound register 0x3dfffff  
 $\Rightarrow$  access violation

(c) Now  $J_4$  is running and it accesses the logical address 0x02f3f12c. Compute the corresponding physical address. Also, check and determine whether this causes access violation.

$J_4$  base register 0x3a000000 + logical address 0x02f3f12c  
 $\Rightarrow$  physical address 0x3cf3f12c < bound register 0x3dfffff  
 $\Rightarrow$  good address, no access violation

6. Consider a 1GB memory system using buddy scheme. The system loads jobs  $J_1, J_2, J_3, J_4$ , and  $J_5$  of size 40MB, 20MB, 160MB, 64MB, and 100MB, respectively.

(a) Compute the starting addresses of each job.

$J_1$  starting at 0x00000000 (0, 64MB), internal fragmentation = 24MB  
 $J_2$  starting at 0x04000000 (64MB, 32MB), internal fragmentation = 12MB  
 $J_3$  starting at 0x10000000 (256MB, 256MB), internal fragmentation = 96MB  
 $J_4$  starting at 0x08000000 (128MB, 64MB), internal fragmentation = 0  
 $J_5$  starting at 0x20000000 (512MB, 128MB), internal fragmentation = 28MB

(b) What is the total internal fragmentation?

Total fragmentation = 24+12+96+28 MB = 160MB

7. Consider a 256MB memory system using simple paging scheme. Each page is of size 16KB. The memory is byte addressable. The page table for Process  $P$  is as follows. Process  $P$ 's address space is only 5 pages. The numbers in the table are all in decimal.

|   |    |
|---|----|
| 0 | 11 |
| 1 | 25 |
| 2 | 40 |
| 3 | 3  |
| 4 | 13 |

- (a) How many bits are required to address the entire memory?  
 (b) How many bits are required to address the offset within each page?  
 (c) What is the maximal possible internal fragmentation  $P$  can have?  
 (d) Given a logical address 0x0000ef5b, compute the corresponding physical address. Also, check and determine whether this causes access violation.  
 (e) Given a logical address 0x00017f5b, compute the corresponding physical address. Also, check and determine whether this causes access violation.

- (a) 28 bits  
 (b) 14 bits  
 (c) 16KB - 1

(d) 0x0000ef5b = 1110 1111 0101 1011  $\Rightarrow$  page 3, mapped to frame 3,  
 $\Rightarrow$  physical address is the same as the logical address

(e) 0x00017f5b = 0001 0111 1111 0101 1011  $\Rightarrow$  page 5, does not exist  $\Rightarrow$  access violation

8. Consider a virtual memory system with cache, main memory, and disk. The cache size is 256KB and has an access time 15ns. The cache hit ratio is 90% (which is considered low). The main memory is 256MB and has an access time 150ns. 99.99% of the memory accesses would find the demanded pages in the main memory and 0.01% of them would generate page faults. The disk size is 20GB and has an access time 1ms. Compute the average memory access time.

Average memory access time

$$= 15 \text{ ns} + (1-0.9) * 150 \text{ ns} + (1-0.9)*(1-0.9999) * 10^6 \text{ ns} = 15 \text{ ns} + 15 \text{ ns} + 10 \text{ ns} = 40 \text{ ns}$$

9. Consider a virtual memory system using 32-bit addressing. The physical memory is of size 4GB. Each page is of size 8KB. Assume that each entry of the page table fits in one word. A two-level page table is used and the root level contains 1K entries. A process P has the following page table.

| Root level page table | Level 2 page table<br>start at 0x00000000 | Level 2 page table<br>start at 0x00000100 | Level 2 page table<br>start at 0x00001000 |
|-----------------------|---|---|---|
| 0x00001000            | frame 150                                 | frame 33                                  | frame 160                                 |
| 0x00000100            | frame 21                                  | disk                                      | frame 512                                 |
| null                  | disk                                      | frame 120                                 | .....                                     |
| .....                 | .....                                     | .....                                     |   |

- (a) How many bits are required to address the root page table?  
 (b) How many bits are required to address each second level page table?  
 (c) Convert the logical address 0000 0000 0100 0000 0101 0000 1111 0101 (binary) to physical address.

(a) 10 bits

(b) 9 bits

(c) 0000 0000 0100 0000 0101 0000 1111 0101

It is root table entry 1, secondary page table entry 2, which is mapped to frame 120 = 0x78

Physical address is: 0000 0000 0000 1111 0001 0000 1111 0101 or 0x000f10f5