

- Virtual memory, demand paging
 - Page replacement policies
 - Follow the algorithms for some example cases
 - The pros and cons of the algorithms
 - The relation of replacement policies to locality and working set sizes
 - Additional issues with page replacement
 - Memory page sharing
 - How to do it and the pros and cons
 - Memory protection
 - How does the process address space look like? Why?
 - What are to be protected?
 - Why there would be access violations?
 - Why there would be the buffer overflow problem and how exactly it happens?
 - Given scenarios and decide whether they may happen and whether the protection mechanisms can help
 - Some questions may involve earlier materials in memory management
- Disk
 - Disk basic terminologies, arm, track, sector, their characteristics
 - Disk access time: seek time, rotational delay, transfer time
 - Some cases consider average
 - Disk access time for files
 - Allocation strategy with access time
 - Disk arm scheduling algorithms and their impact in file systems
 - RAID with different designs
 - Basic design and data layout on parallel disks
 - Read/Write behaviors and performance
 - Fault tolerance capability and performance after failure(s)
 - Comparisons between the designs
 - RAID-DP
 - Understand the disk layout and recovery mechanisms
- Clock
 - The design of the clock itself (the registers and their values)
 - OS supported timer, the data structure
 - Why this data structure?
 - How it works, especially for some specific examples
 - Potential pros and cons with an alternate design
- Terminal
 - The basic concept on how keyboard, display, and OS are interacting
 - The intelligent keyboard-monitor interactions
- File system
 - Basic file system concepts
 - Different allocation strategies and their pros and cons
 - Which method is suitable for what type of file accesses, may give scenarios to ask you to rank various methods
 - Basic Unix file system

- Inode structure, may change it to test your understanding
 - Directory structure and how to find the inum of a file
 - From inum, how to find the inode of the file
 - From inode, how to find a certain data block of the file
 - If we want to access a certain block of the file, where is it on the disk (track number, sector number, etc.)
- The cylinder groups in FFS, its design goal, and detailed design
 - Same as the key points as above, but associate with cylinder group
- Journaling and fault tolerance in FS
 - Basic hardware support for dealing with failures
 - Given a failure scenario, what would be the consequence for the file system? E.g., file block lost, wrong data or garbage in the block, etc.
 - If some problem is detected, what activities have occurred to cause the problem? What actions can be taken to recover from the problem?
- LFS
 - Basic concept of the approach
 - inode block problem and the imap solution
 - Layout of the information in a segment and why we need each part
 - Final solution: imap in CR, why it has to be there in CR? Why the solution can support delayed write?
 - How to perform read?
- NFS
 - The concept of v-node and its relation to i-node
 - The basic flows regarding how to mount an external disk to a local directory (or drive) and how to retrieve an NFS file, what data structures need to be maintained
 - The data structure for maintaining file descriptors of a process
- File system design concepts
 - Compare different file system designs in general and for some specific scenarios
 - Given a different FS design or alter an existing design, what are the pros and cons or even flaws
- Processor scheduling
 - Basic processor scheduling policies
 - Follow the algorithms to schedule a given set of tasks
 - MLFQ and Unix MLFQ are especially important
 - Potential problems, pros and cons, comparisons
 - Compare the algorithms for some given scenarios
 - Real time scheduling algorithms
 - Follow the algorithms to schedule a given set of tasks
 - Compare the algorithms for some given scenarios
 - Fully understand priority inversion and resolution
 - May give modified algorithms
 - Understand their pros and cons and compare them with existing algorithms
 - Some of the questions may involve the earlier materials on process states

- Project related questions
 - Fork, thread, pipe, signal, socket
 - The basic concept of these functions
 - Under an example scenario, which ones are the best to use and “why”
 - May need to use multiple functions
 - There will be no question on thread specific behaviors, but combined with other functions
 - The behaviors of a program with these functions
 - Try out the behaviors if these functions are called differently
 - Projects 2 and 3
 - Basic designs
 - Thread synchronization and semaphore coding for Project 3
 - Different handling between the ready queue and terminal queue (how they are different and why)
 - Socket communication steps and select command

- Sample Questions
 - Variations of questions in the homework assignments
 - Concept based questions
 - Which of the following page replacement algorithms is most similar to the least recently used algorithm?
 - Why is there the buffer overflow problem when any virtual address is supposed to be valid?
 - Rank the following processor scheduling algorithms in terms of the response time they yield
 - Rank the following parallel disk designs in terms of their fault tolerance levels
 - Most likely this type of questions will involve algorithms that are modified from what you have learnt
 - More potential questions have been discussed at a high level in the review, you need to study in more depth according to the review statements
 - The sample questions are given to show the type of questions you may get besides those appear in homework #3, but the exam questions will be different