

Virtual Memory Management

Memory Management

❖ Previous solutions

- Need to know the memory requirement for the process
 - Fixed partitioning, dynamic partitioning, simple paging
 - Not easy to do so
- Need to allocate a contiguous memory space for a process
 - Fixed partitioning, dynamic partitioning
 - So that addressing on the fly can be done efficiently
 - But has fragmentation problem
 - Even there is sufficient memory, allocation may not be possible
- ⇒ Paging is better
 - Simple paging requires a process to know #pages it uses in advance
 - ⇒ But not difficult to allocate new pages when they are needed
 - How about needing more pages the memory can offer?
 - How about page table overhead

Virtual Memory

❖ What is the size limit of a process

- A process can be as large as it can address
 - Can be much larger than the physical memory size
- The limit is set by the address space
 - 32 bits address → 4G address space
 - Even if it is big enough for each process, it is too small to address modern physical memory (beyond 4GB)
 - 64 bit address space → 16 exa bytes
- Not all parts of the address space of a process have to be in memory at the same time
 - For regular programs ⇒ A lot of the address space is empty
 - For huge programs ⇒ Only those parts that are needed during current and near future execution need to be in memory
- The “memory” for each process is “virtual”

Virtual Memory and Demand Paging

❖ Virtual Memory

- Virtual addresses are used for memory accesses
 - Virtual addresses are mapped to physical addresses on the fly
 - The same in all memory management schemes
- The size of each process is as large as the address space
 - Only true in virtual memory paradigm

❖ Demand paging

- Address space is divided into pages
- A page is brought into physical memory only when it is needed ⇒ This is called demand paging
- But where are the remaining pages? ⇒ Disk of course

Virtual Memory and Demand Paging

❖ Demand paging

➤ Swap space

- All pages (that contain something) are stored on the disk
- In the swap space allocated specifically for this purpose
 - Generally OS allow admin to set up the swap space size

➤ If a needed page is not in memory

- A **page fault** occurs
- OS brings the page from swap space into memory
- ...

❖ ⇒ Memory hierarchy for memory content

- At least need disk to realize virtual memory, demand paging
- Cache for improving the access performance
 - For page table as well as content

Memory Hierarchy

❖ Register

❖ Cache

❖ Memory

❖ Disk (swap space, not files)

- Many new research consider flash memory as another layer

❖ Down the hierarchy

- Decreasing cost
- Increasing capacity and access time

Performance: Memory Hierarchy


❖ Memory hierarchy characteristics

- Register: ~ 1 ns or less
- Cache access time: ~ 1 ns to 10 ns
 - (for i7) L1 cache: ~4 cycles, L2: ~12 cycles, L3: ~40 cycles
- Memory access time: ~ 50 ns
 - (overall time, not just the memory CAS latency, which is ~10ns)
- Disk access time: ~ 0.1 ms
- Cache hit rate
 - Probability that a word is in cache
- Memory hit rate
 - Probability that a word is not in cache but in memory

Performance: Memory Hierarchy

❖ Memory access

- Consider cache, memory, disk
- Frequently used pages are stored in the cache
- For a memory access



Memory Hierarchy (3 Levels)

❖ Average memory access latency

$$\begin{aligned} & 0.95 * 10 \text{ ns} + \\ & [(1 - 0.95) * 0.999] * (100+10) \text{ ns} + \\ & [1 - (0.95 + (1 - 0.95) * 0.999)] * (100000+100+10) \text{ ns} \\ & = (9.5 + \sim 5.5 + \sim 5) \text{ ns} \qquad = 1 * 10\text{ns} + .05 * 100\text{ns} + .00005 * 100000\text{ns} \\ & = 20 \text{ ns} \end{aligned}$$

Cache
Access time = 10 ns
Hit rate = 95%

Memory
Access time = 100 ns
Hit rate = 99.9%

Disk
Access time = 100μs

Virtual Memory Issues

❖ Allocation

- Different issue in paging schemes, not about where, but ...
- How many pages should be loaded for each process

❖ Addressing

- How to map a logical address to a physical address?
- How to know which frame each page is mapped to?
- Page table
 - Need a page table to keep track of the pages of a process
 - In cache, or memory, or disk
 - How to store the page table?
 - How to efficiently search the page table?

Virtual Memory Issues

❖ Replacement

- A page that has been brought into memory may reside there till it is replaced
- When a new page is to be brought into memory and there is no free frames
 - Need to replace an existing page
 - How to choose the process and the page to replace?

❖ Working set management

- Working set: the set of pages required during execution
- What is the best working set size?

Virtual Memory Issues

❖ Protection

- Protect the address space of one process from being accessed by another process

❖ Performance metrics

- No longer consider fragmentation
- Time for each memory access \approx Time for page table lookup
- Number of page faults in a time unit

Addressing: Page Table

❖ Where to put the page table

- In main memory, at least
 - Cache is too small, PT on disk will yield ridiculous performance
- How many pages each process has?
 - 32-bit system: each process has $4\text{GB}/8\text{KB} = 0.5\text{M}$ pages
 - 64-bit system?
- \Rightarrow Page table only stores the pages that are not empty
 - What should the data structure for the page table be so that search can be done efficiently? \Rightarrow Need a good design!!!
 - Array with binary search? Hash table with linear probing?
 - Potentially more than 1 extra references to M for each M access
 - No matter how good the data structure is, it will be at least one extra memory reference
- \Rightarrow Memory hierarchy for page table

Page Table for Virtual Memory

❖ Page table designs

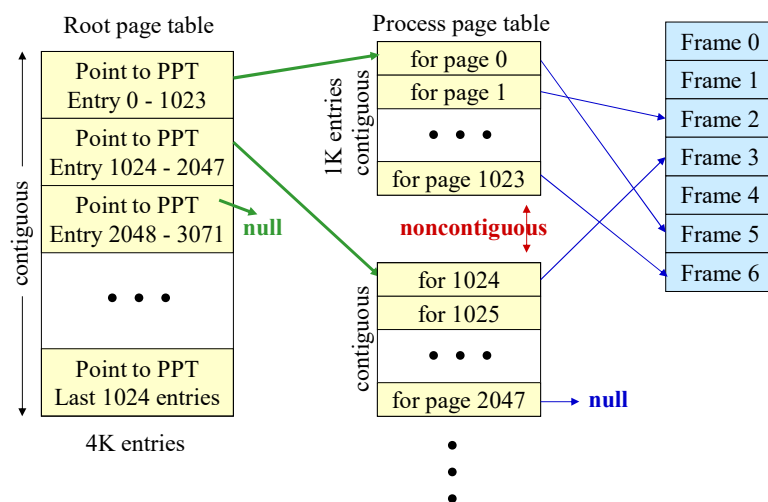
- Process page table
 - Full indexing for each process
 - Can map to memory and disk
- Inverted Page Table (IPT)
 - Maps memory frames to pages in processes
 - Only address pages in memory
- Translation lookaside buffer (TLB)
 - Only address pages in memory

Two-Level Process Page Table

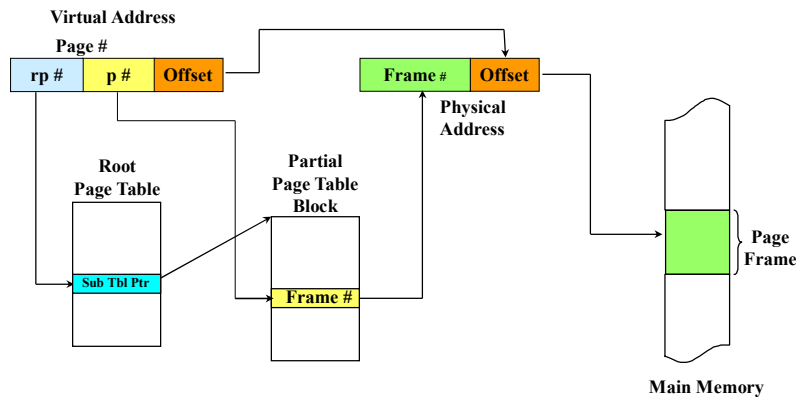
❖ How to search process page table efficiently

- Contiguous page table
 - 4GB virtual address space (32 bits), assume 4KB page size
 - Page table has 1M entries for **each process**
- Most of the entries in the table are empty
- Non-contiguous
 - How to search? Efficiency will be a problem
- Two level page table
 - Use another table to index the page table (**root page table**)
 - The low level ones are the **process page tables**
 - E.g., 1K entries in the RPT
 - Each entry in RPT points to the nonempty PPT block

Two-Level Page Table (for one process)



Two-Level Page Table (for one process)



Addressing in Two-Level PT

❖ Addressing

- 4GB address space, 1KB per page \Rightarrow 4M pages
- Root page table has 16K entries
 - Use 14 bits for addressing the root page table
- Each partial page table block is 256 entries
 - Use 8 bits for addressing the page table block
- **0000000000001000001010000101011**
 - Root table page number = 10 = 2
 - Partial Page block index = 1010 = 10
 - Page offset = 101011
- **00000000000000000101010000101011**
 - Frame number = 21 = 10101

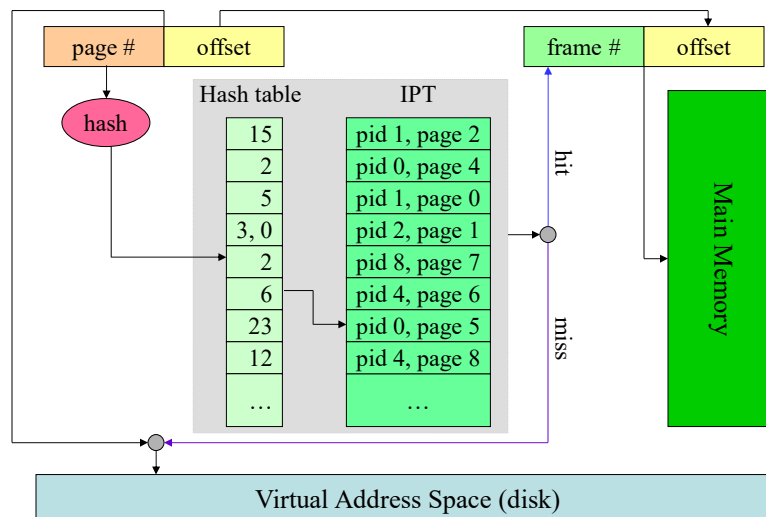
0		
1		
2		
...		
9		
10		

0	7	P.512
1	40	P.513
2	1	P.514
...		
9	15	P.521
10	21	P.522
...		

Inverted Page Table (IPT)

- ❖ Inverted: Map frame # to processe ID/page #
 - # IPT entries = # frames in physical memory
 - Indexed by physical page frames
 - Useful when a frame is modified
 - Easy to locate the corresponding process ID/page #
- ❖ Find page # from frame # can be very expensive
 - Use hash table to map page # to frame #
 - Require at least two extra memory accesses
 - One for hash table, one for IPT accesses
 - Collision may occur in the hash table
 - Need an additional field for chaining

Addressing in IPT



Addressing in IPT

❖ Addressing

- Given (process#, page#) tuple
- Locate the corresponding physical frame

❖ hash function $h = \text{process\#} \otimes \text{page\#}$

❖ Example 1: Find frame number for (0,5)

- $h(0, 5) = 0000 \otimes 0101 = 0101 = 5$
- $\text{hashtable}[5] = 6$
- $\text{IPT}[6] = (0,5) \Rightarrow$ Got the match
- (process 0, page 5) is in memory frame 6

Addressing in IPT

❖ Example 2: Find frame number for (1, 2)

- $h(1, 2) = 0001 \otimes 0010 = 0011 = 3$
- $\text{hashtable}[3] = 3, 0$
- $\text{IPT}[3] = (2,1) \Rightarrow$ does not match
- $\text{IPT}[0] = (1,2) \Rightarrow$ Got the match
- (process 1, page 2) is in memory frame 0

Addressing in IPT

❖ Example 3: Find frame number for (8, 9)

- $h(8, 9) = 1000 \otimes 1001 = 0001 = 1$
- $\text{hashtable}[1] = 2$
- $\text{IPT}[2] = (1,0) \Rightarrow$ does not match
- (process 8, page 9) is not in memory
- Issue a page fault

Translation Lookaside Buffer

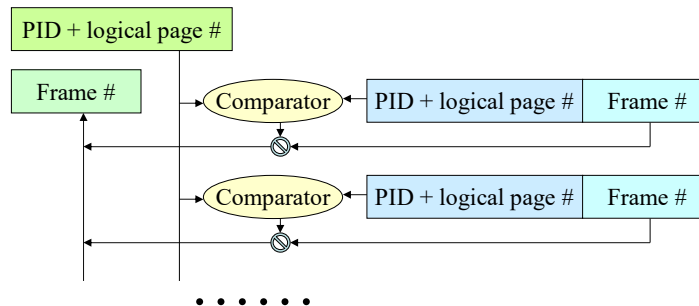
❖ For every memory access

- Need to search the page table in memory and map the logical page number to physical frame number
- \Rightarrow Too expensive
- Could cache the page table
 - Still relatively expensive
 - How to index the selected pages in the cache (non-contiguous \Rightarrow cannot be indexed \Rightarrow need to search sequentially)
- Solution
 - Store the page-to-frame mappings in a fast cache
 - Use associative memory to enable single cycle parallel search
 - This is called the translation lookaside buffer (TLB)

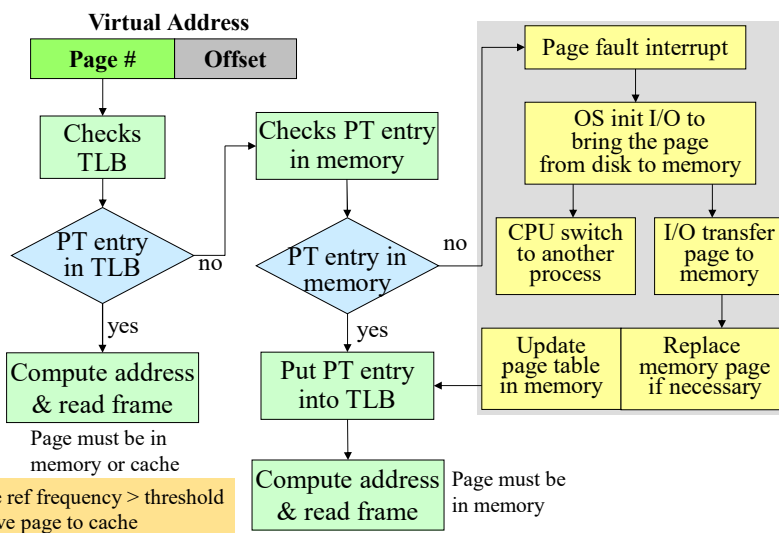
Translation Lookaside Buffer

❖ TLB

- Use process ID and logical page number as the key
- Search for the entry in parallel to get the frame number
- Latency ~ 5-10 ns (not really register speed, ~ cache speed)
- The page currently referenced is definitely in TLB, for locality



Addressing -- Obtain Frame Number



Working Set Management

- ❖ Working set (resident set)
 - The set of pages that a process needs within a time interval
- ❖ Working set size for each process
 - Depend on the time interval considered
 - If it is small
 - Higher level of multiprogramming
 - Better for time-sharing systems
 - May result in **thrashing**
 - Page faults occur every few instructions
 - If it is too big
 - May have lower level of concurrency

Working Set Management

```
int A[128,128], B[128,128];
int i, j, x, y;
for (i=0; i < 128; i++)
    for (j=0; j < 128; j++)
        { B[i,j] = A[i,j] + x + y;
          x = (x * i) % 128; y = (y * j) % 128; }
```

- ❖ Memory
 - Page size = 1K bytes; size(int) = 1 word = 4 bytes
 - Program and x , y , i , and j are stored in one page
 - $A[0,0]$ - $A[0,127]$, $A[1,0]$ - $A[1,127]$ in a page
 $B[0,0]$ - $B[0,127]$, $B[1,0]$ - $B[1,127]$ in a page

Working Set Management

❖ When to load the working set

- Load on demand
 - A page will be loaded when needed
 - Suitable for new jobs
 - May take some initialization time for a process to get all its needed pages
- After being swapped out and swapped in
 - Load back pages on demand will incur a big overhead
 - Preload the working set before process starts to run

❖ Variable working-set size

- More page faults \Rightarrow Larger working set size
- Pages that are not used for a time period \Rightarrow Remove the page and reduce the working set size

Page Replacement Policies

❖ Principle of locality

- A program that references a location l at some point of time is likely to reference the same location l and locations in the immediate vicinity of l in the near future
 - 90% of the execution time is spent on loops
- Most of the time, a program is executed sequentially
- A program tends to favor a subset of its pages during a time interval

Page Replacement Policies

❖ Replacement policies

- When a working set of a process is full and a new page is needed \Rightarrow replace an existing page
- Which page to replace

❖ Global versus process based

- Replacement for each process
 - Each process has a working set size
- Global
 - There is no need to consider working set size for each process
 - May replace any page in the memory

Page Replacement Policies

❖ Replacement policies

- When a working set of a process is full and a new page is needed \Rightarrow replace an existing page
- Which page to replace

❖ Policies

- First-in-first-out policy (FIFO)
- Least recently used policy (LRU)
- Clock policy
- Modified clock policy
- Aging policy

Page Replacement Policies

❖ First-in-first-out policy (FIFO)

- Simple (use a pointer points to the oldest page)
- The oldest page may be used more recently

❖ Least recently used policy (LRU)

- Replace the least recently used page
- Comply with the principle of locality
- High implementation overhead
 - Hard to keep track of all references
 - Use a stack for each process

Clock Policy

❖ Pointer *mrp*

- points to the most recently **loaded** page

❖ Flag *used*

- indicates whether a page has been used recently

❖ Page replacement

- Scan and find the **first page with *used* = 0**
 - *mrp* points to the most recently loaded page
 - Scan from the page that is after where *mrp* is pointing to
- During scanning, **set *used* → 0** if it was 1
- If no page with *used* = 0 in the first round → will definitely get one during second round

Modified Clock Policy

- ❖ Same as before: *mrp* and *used*
- ❖ Flag *dirty*
 - Indicate whether a page has been modified
 - *dirty*=1 → the page have to be written out before clear the flag to 0
- ❖ Principle
 - Same as clock
 - First try to find a *used* = 0 page
 - Among them, *dirty* = 0 should be considered first

Modified Clock Policy

- ❖ Page replacement
 - Scan from *mrp*
 - First scan
 - Find the **first page with $used = 0 \wedge dirty = 0$**
 - No change to the *used* and *dirty* bits
 - Second scan, if no suitable pages found
 - Find the **first page with $used = 0 \wedge dirty = 1$**
 - During scan, **set $used \rightarrow 0$** if it was 1
 - Never change dirty bit
 - Third scan, if still no suitable pages found
 - Find the **first page with $used = 0 \wedge dirty = 0$**
 - Fourth scan, if still no suitable pages found
 - Find the **first page with $used = 0 \wedge dirty = 1$**

Will definitely
find one now

Example for Page Replacement Policies

- ❖ Page use pattern: 0 1 5 0 1 4 0 1 2 0 2 3 0 6 3
- ❖ Working set size: 3
- ❖ FIFO (10 page faults)

0 1 5 0	0 1 2 0
0 1 5 1	0 1 2 2
0 1 5 4	0 1 2 3
1 5 4 0	1 2 3 0
5 4 0 1	2 3 0 6
4 0 1 2	3 0 6 3
0 1 2 0	3 0 6

Example for Page Replacement Policies

- ❖ Page use pattern: 0 1 5 0 1 4 0 1 2 0 2 3 0 6 3
- ❖ LRU (7 page faults)
 - Minimal page faults required: 7

0 1 5 0	0 1 2 0
1 5 0 1	1 2 0 2
5 0 1 4	1 0 2 3
0 1 4 0	0 2 3 0
1 4 0 1	2 3 0 6
4 0 1 2	3 0 6 3
0 1 2 0	0 6 3

Least recently used ← — Most recently used

Example for Page Replacement Policies

❖ Page use pattern: 0 1 5 0 1 4 0 1 2 0 2 3 0 6 3

❖ Clock (9 page faults)

➤ * represents used bit, ↓ next to the last loaded pages

0* ↓ 1	4* ↓ 1 5 0	2* ↓ 0* 1 3
0* 1* ↓ 5	4* 0* ↓ 5 1	2* 0 3* 0
0* 1* 5* 0	4* 0* 1* 2	2* 0* 3* 6
0* 1* 5* 1	2* ↓ 0 1 0	6* ↓ 0 3 3
0* 1* 5* 4	2* ↓ 0* 1 2	6* ↓ 0 3*

Example for Page Replacement Policies

❖ Page use pattern: 0 1 5 0 1 4 0 1 2 0 2 3 0 6 3

❖ Clock (9 page faults)

➤ * represent

Scan and find page with (used = 0)
While scanning, change (used → 0)

⇒ 0 1* 5* | 4

⇒ 0 1 5* | 4

⇒ 0 1 5 | 4

Finish one round, no page selected

Start again, page 0 selected

⇒ 4* 1 5 |

0* ↓ 1		
0* 1* ↓ 5		
0* 1* 5* 0		
0* 1* 5* 1		
0* 1* 5* 4		
	4* ↓ 1 5	
	2* ↓ 0* 1 2	6* ↓ 0 3 3
		6* ↓ 0 3*

Example for Page Replacement Policies

Scan and find page with (used = 0)

Page 1 selected

⇒ 4* 0* 5 |

1 4 0 1 2 0 2 3 0 6 3

➤ * represents used bit, ↓ next to the last loaded pages

0* ↓ 1	4* ↓ 1 5 0	2* ↓ 0* 1 3
0* 1* ↓ 5	4* 0* ↓ 1	2* 0 3* 0
0* 1* 5* 0	4* 0* 1* 2	2* 0* 3* 6
0* 1* 5* 1	2* ↓ 0 1 0	6* ↓ 0 3 3
0* 1* 5* 4	2* ↓ 0* 1 2	6* ↓ 0 3*

Example for Page Replacement Policies

❖ Page use pattern: 0 1 5 0 1 4 0 1 2 0 2 3 0 6 3

❖ Clock (9 page faults)

➤ * represents used bit, ↓ next to the last loaded pages

Scan and find page with (used = 0)

While scanning, change (used → 0)

⇒ 4 0 1 | 2

Finish one round, no page selected

Start again, page 0 selected

⇒ 2* 0 1 |

0* ↓ 1	4* ↓ 1 5 0	
0* 1* ↓ 5	4* 0* ↓ 1	
0* 1* 5* 0	4* 0* 1* 2	2* 0* 3* 6
0* 1* 5* 1	2* ↓ 0 1 0	6* ↓ 0 3 3
0* 1* 5* 4	2* ↓ 0* 1 2	6* ↓ 0 3*

Example for Page Replacement Policies

❖ Page use pattern: 0 1 5

❖ Clock (9 page faults)

➤ * represents used bit,

Scan and find page with (used = 0)
While scanning, change (used → 0)
⇒ 2* 0 1 | 3
Page 1 selected
⇒ 2* 0 3* |

0* ↓ 1	4* ↓ 1 5 0	2* ↓ 0* 1 3
0* 1* ↓ 5	4* 0* ↓ 5 1	2* 0 3* 0
0* 1* 5* ↓ 0	4* 0* 1* ↓ 2	2* 0* 3* ↓ 6
0* 1* 5* ↓ 1	2* ↓ 0 1 0	6* ↓ 0 3 3
0* 1* 5* ↓ 4	2* ↓ 0* 1 2	6* ↓ 0 3*

Example for Page Replacement Policies

❖ Page use/write pattern

0 1w 5 0 1w 4w 0 2 0 1w 2w 4

❖ Modified Clock Policy (x represents modified)

0* ↓ 1w	0* 1*x 5* ↓ 4w	4x 2* ↓ 0* 1w
0* 1*x ↓ 5	4*x ↓ 1x 5 0	1*x ↓ 2* 0 2w
0* 1*x 5* ↓ 0	4*x ↓ 1x 0* 2	1*x ↓ 2*x 0 4
0* 1*x 5* ↓ 1w	4x 2* ↓ 0* 0	1*x ↓ 2*x 4*

Example for Page Replacement Policies

❖ Page

❖ Modified

Scan and find page with (used=0, dirty=0)

⇒ 0* 1*x 5* | 4w → no page selected

2nd scan, find page with (used=0, dirty=1)

While scanning, change used bit

⇒ 0 1x 5 | 4w → still no page selected

Scan and find page with (used=0, dirty=0)

Page 0 selected

⇒ 4*x 1x 5 |

0*	↓	1w	0*	↓	1*x 5*	4w	4x	↓	2*	↓	0*	1w
0*	↓	1*x	↓	5	4*x	↓	1x	↓	5	↓	0	2w
0*	↓	1*x	↓	5*	↓	0	4*x	↓	1x	↓	0*	2
0*	↓	1*x	↓	5*	↓	0	4x	↓	2*	↓	0*	0

Example for Page Replacement Policies

❖ Page use

❖ Modified

Scan and find page with (used=0, dirty=0)

Page 5 selected

⇒ 4*x 1x 0* |

0*	↓	1w	0*	↓	1*x 5*	4w	4x	↓	2*	↓	0*	1w
0*	↓	1*x	↓	5	4*x	↓	1x	↓	5	↓	0	2w
0*	↓	1*x	↓	5*	↓	0	4*x	↓	1x	↓	0*	2
0*	↓	1*x	↓	5*	↓	0	4x	↓	2*	↓	0*	0

Example for Page Replacement Policies

❖ Page use/wr

0 1w 5

❖ Modified C

Scan and find page with (used=0, dirty=0)


⇒ 4*x 1x 0* | 2 → no page selected

2nd scan, find page with (used=0, dirty=1)

While scanning, change used bit

⇒ 4x 1x 0* | 2 → page 1 selected

⇒ 4x 2* 0* | → write page 1 to disk



Example for Page Replacement Policies

❖ Page use/write pattern

0 1w 5 0 1w 4w

❖ Modified Clock Policy

Scan and find page with (used=0, dirty=0)

⇒ 4x 2* 0* | 1w → no page selected


2nd scan, find page with (used=0, dirty=1)

While scanning, change used bit

⇒ 4x 2* 0 | 1w → page 4 selected

Write page 4 back to disk, replace it

⇒ 1*x 2* 0 |



Aging Policy

❖ Aging policy (global page replacement)

- Increase used field whenever it is accessed
 - Set the leftmost bit of the age to 1
- System periodically scan the memory and reduce the “used” value
 - Right shift the age value
- A page is to be removed if it reaches 0
 - Does not need to actually remove till receiving a replacement request
 - Replace a page with the lowest age number
- Special hardware
 - Special memory for the age vector, able to shift, set, select min

Aging Policy

❖ Aging policy

- Example
 - 3 pages in memory, page use pattern: 0 1 5 0 1 4 0 1 2 0 2 3 0 6 3
 - 8-bit age vector, scan every 1 sec (each sec has 1 page accesses)
- Result
 - Start: 00000000 (-) for all frames
 - Access 0: 10000000 (0), 00000000 (-), 00000000 (-)
 - Access 1: 01000000 (0), 10000000 (1), 00000000 (-)
 - Access 5: 00100000 (0), 01000000 (1), 10000000 (5)
 - Access 0: 10010000 (0), 00100000 (1), 01000000 (5)
 - Access 1: 01001000 (0), 10010000 (1), 00100000 (5)
 - Access 4: 00100100 (0), 01001000 (1), 10000000 (4)
 - Choose the page with the lowest age to replace

Other Page Replacement Considerations

❖ Global page replacement

- Consider all processes together
- E.g., Linux page replacement policy
 - Aging policy with 8-bit aging vector

❖ Load Control

- How many processes should reside in memory
 - In global policy, this cannot be controlled by working set size
- Control the level of multiprogramming
- Too many processes → high frequency of page faults
- Too few processes → reduced level of concurrency

Other Issues in VM

❖ Frame Locking

- Lock pages that cannot be replaced
 - e.g. OS code and data
- When a page is just brought in and before it is used
 - Process A is switched out when having a page fault
 - When the page is brought in, A may not be scheduled to run

Other Issues in VM

❖ Page sharing

- In many situations, multiple address spaces may have exactly the same pages
 - E.g., library functions, kernel functions
 - Only one copy is needed in physical memory
 - Page tables of all the processes for the shared page point to the same memory frame
- Copy on write
 - When a process writes to the shared page, then the shared page is copied and a private copy is given to the process
 - Save physical memory space as well as copying time

Other Issues in VM

❖ Page sharing

- Fork: child processes initially have exactly the same copy as the parent process
- No need to do actual copy unless there is a write
 - What needs to be copied?
- Frequently exec follows fork, and all the copying after fork is wasted
- With copy on write, the saving can be significant

Other Issues in VM

❖ Protection

➤ How to ensure that each process accesses its own memory


- Is it a problem in virtual memory?
- Why there are access violations?

➤ Memory layout

- Stack: call stack
- Heap: dynamic references
 - E.g., those allocated by malloc
- OS: in the address space of every process
 - No context switch for kernel mode
 - Share the same frames by all processes
 - Use copy on write when data differs

➤ Frame metadata

- Accessed, modified, valid, read only, mode (user/kernel)




Other Issues in VM

❖ Protection

➤ Is the entire address space mapped in page table?

- The empty region between stack and heap is not
 - No page for it (neither in memory nor in swap space)
- OS allocate a certain size of memory for stack
 - Can be changed, but generally are fixed
- Heap is requested page by page
 - When malloc, if no space in existing heap page, system call sbrk/mmap is invoked to increase the heap space, generally it is page aligned
 - A marker marks the boundary of the heap space sbrk/mmap moves it
 - sbrk/mmap can cause a new entry (or new entries) being inserted to the page table




Other Issues in VM

❖ Protection

➤ Why there are access violations?

- If accessing OS region
 - Page table points to shared OS frames
 - The mode bit for these frames are set to “kernel mode”
 - These OS frames can only be accessed in kernel model
 - Set the mode register to kernel mode
 - Only allowed in system calls, exception handlers, etc.
- If accessing the empty region
- If accessing address 0
 - Address 0 is frequently reserved specifically in the address space for null
- Even though the above are in the address space of the process!!!

Why put OS code in user space?
⇒ Avoid context switch
⇒ Need to switch to kernel mode



Other Issues in VM


❖ Protection

➤ Call stack

- When there is a procedure call, the in/out parameters, the local data, and a return address are put into stack

➤ Stack buffer overflow

- A procedure with a local buffer is put into stack
- Caller provides input longer than the buffer size
- May put malicious code in the buffer, longer input overwrites the return address, causing branching into the malicious code
- Can attack the system, but only if the process being attacked by buffer overflow has the privilege



Other Issues in VM

❖ Protection

➤ Zeroing pages

- When a page is brought into memory for a process, the entire frame is zeroed to ensure that one process does not get to read another process's previous memory content
 - Because zeroing is expensive, it can be done in a lazy way

Readings

❖ Sections 1.5, 1.6

❖ Sections 8.1, 8.2