

Memory Management

Memory Addressing

- ❖ `main()`
 { static int a, b, c;
 a = 5; b = 3; c = a + b;
 printf ("addresses: %d, %d, %d\n", &a, &b, &c);
 }
- ❖ Question: Run two instances of the same program, will the addresses of a be the same?

System Operations

- ❖ Compiler uses offset for addressing
 - Need to do this because program may reside in different memory locations even during execution
- ❖ Linker combines all object files into an executable - unifies the address offset
- ❖ Loader loads executable into memory, the base address is determined after loading

Memory Allocation

- ❖ Important elements in memory allocation
- ❖ Allocation Policy
 - Where in the memory to allocate a program
- ❖ Relocation
 - Allow the program being swapped out and swapped back in at a different memory location
- ❖ Addressing
 - Has to be offset, otherwise, the instructions has to be changed continuously
 - How to compute physical addresses from offsets
 - Computation needs to be very efficient

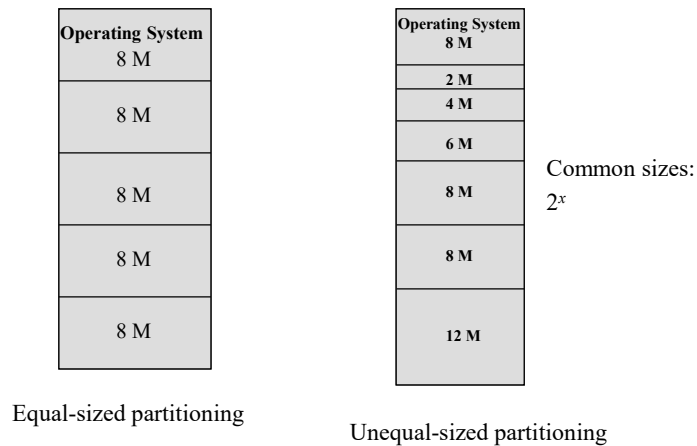
Memory Allocation

- ❖ Important elements in memory allocation
- ❖ Protection
 - A process should only access its designated memory locations
- ❖ Performance Metrics of memory allocation algorithms
 - Fragmentation
 - External fragmentation
 - Internal fragmentation
 - Allocation time
 - Time for executing the allocation algorithm

Allocation Policies

- ❖ Fixed Partitioning
- ❖ Dynamic Partitioning
- ❖ Simple Paging
- ❖ Simple Segmentation
- ❖ Virtual Memory and Demand Paging
- ❖ Virtual Memory and Demand Segmentation

Fixed Partitioning



Fixed Partitioning

- ❖ Partitions are set up at system initialization time
 - Programs are still allocated dynamically
 - A program can only be in one partition
- ❖ Advantages
 - Simple
- ❖ Disadvantages
 - Internal fragmentation
 - Cannot handle processes larger than the biggest partition but smaller than the entire memory

Fixed Partitioning

❖ Allocation policy

- For equal-sized partitioning → just allocate any available partition
 - Very simple
 - Small process ⇒ big fragmentation, Big process ⇒ No slot can fit
- For unequal-sized partitioning →
 - A job goes to the best-fitted partition that is currently available
 - Possible that the best-fit partition is in use
 - Internal fragmentation
 - Each job wait for the best-fitted partition
 - Less fragmentation
 - If many jobs fit one particular partition, these jobs have to wait even though there are many other available partitions

Dynamic Partitioning

❖ Fixed partitioning is not good

- Internal fragmentation
- Cannot handle large processes


❖ Dynamic partitioning

- Memory is not partitioned in the beginning
- One whole partition at initialization time

❖ Allocation policy

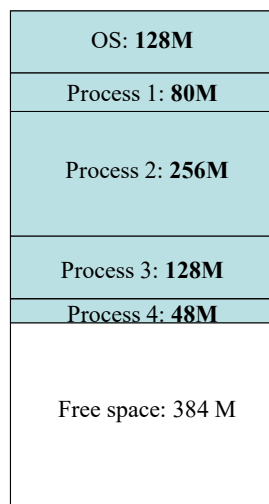

- Allocation size = Process size
- No internal fragmentation

Dynamic Partitioning



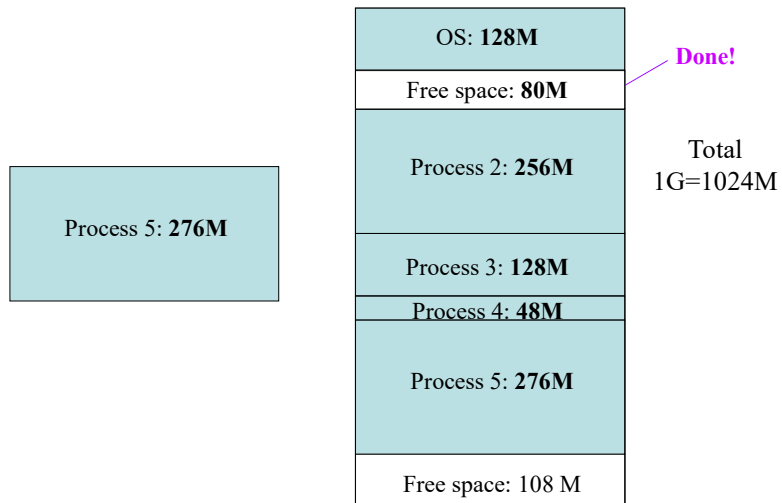
Total
1G=1024M

Dynamic Partitioning

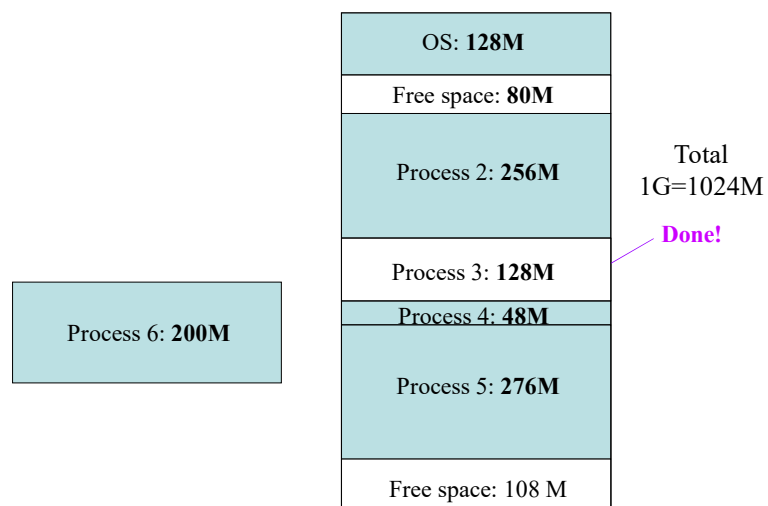


Total
1G=1024M

Dynamic Partitioning



Dynamic Partitioning



Dynamic Partitioning

Cannot allocate


Even though memory has sufficient space left

(External fragmentation!!!)

(Memory may be overly fragmented after a while)

Need to compact memory

Process 6: **200M**



Total
1G=1024M


Dynamic Partitioning

❖ Allocation Algorithms

– Where should a new process be allocated?


- First fit
 - Find first slot which is large enough for the new process
- Best fit
 - Select the slot with size closest to the requested size
- Worst fit
 - Select the slot with the largest size
- Next fit
 - Start to search from the most recently allocated slot
 - Find the first fitting slot
- Buddy Scheme
 - Split a partition into sizes of 2^x and allocate
 - Merge partitions (if possible) after a process leaves

Allocation Algorithms - DP




Allocation Algorithms - DP

Buddy




Allocation Algorithms - DP

Buddy



Allocation Algorithms - DP

Buddy



Allocation Algorithms - DP

❖ Buddy scheme

- Has both internal and external fragmentation problems
- Easy table look up (maintain a tree)

❖ First, best, worst, next fit

- All has external fragmentation problem
- Performance studies show they perform similarly
- Need to keep a free list
 - Best fit and worst fit: list sorted in the order of free partition size
 - First fit and next fit: list sorted in the order of starting address
- When a job leaves → return to free list → merge free slots
 - Memory manager checks the neighboring slots
 - Found a pair of free neighboring partitions ⇒ Merge them
 - If free list is sorted by size ⇒ Hard to locate neighbors

Addressing/Protection - DP

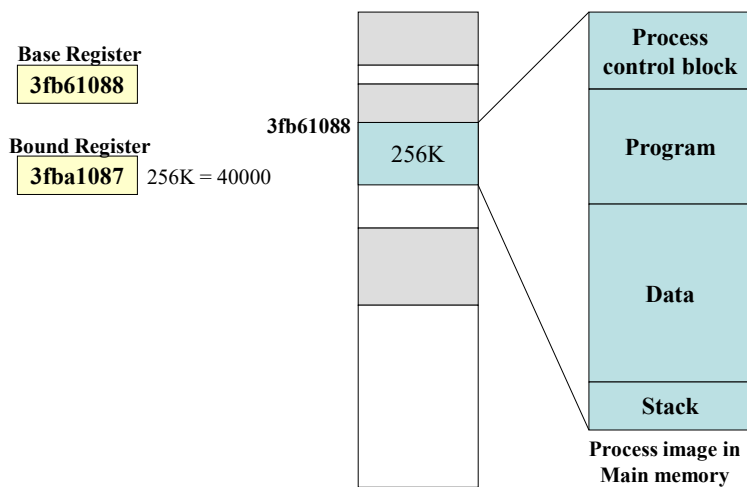
❖ Addressing

- physical address = base address + relative address
- (relative address, offset, logical address) are the same
- (physical address, absolute address) are the same
- Use a base address register to store the base address

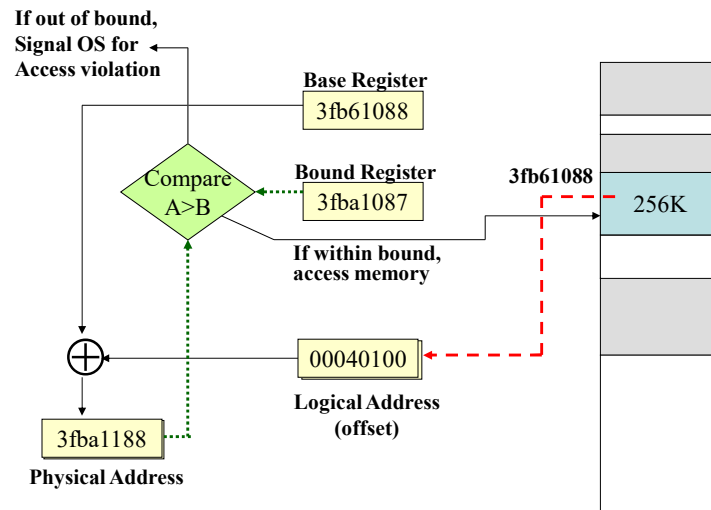
❖ Protection

- physical address limit = base address + program size
- Use a bound register to protect against access violations

Addressing/Protection - DP



Addressing/Protection - DP



Simple Paging

❖ Dynamic partitioning

- Major fragmentation problem
 - Have sufficient space but do not have sufficient consecutive space
- Better solution
 - Let a program to occupy non-consecutive regions
 - Use a table to keep track of the regions used for a program
 - But table may be too long
 - Use page as a unit → reduce number of table entries

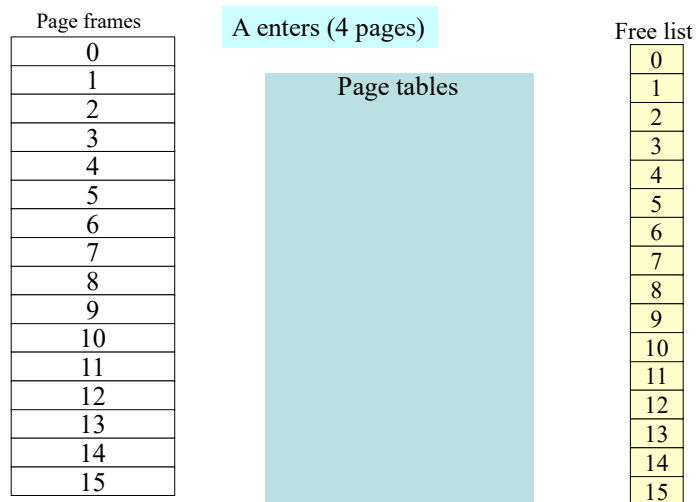
⇒ Paging scheme

- Memory is divided into fixed-size blocks
- Physical memory block → **page frames**
- Process address space block → **pages**

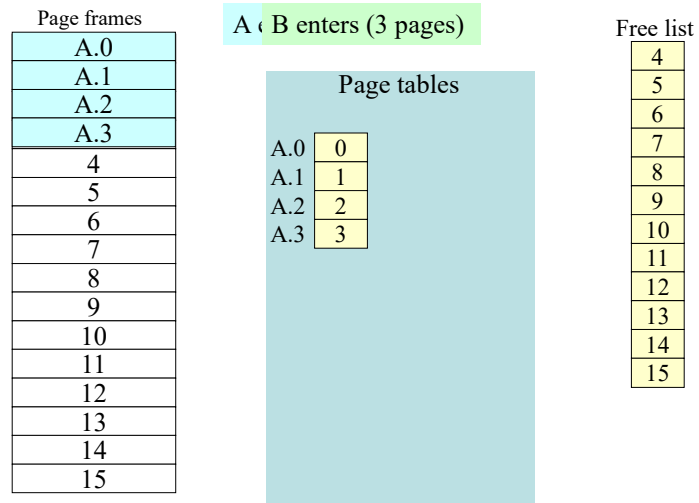
Simple Paging

- ❖ OS maintains a list of free frames
- ❖ Each process maintains a page table in PCB
- ❖ Process does not need to have contiguous frames
- ❖ Page size
 - Typically, 1KB to 8KB
 - Too big → internal fragmentation
 - Too small → large page table
- ❖ Performance
 - little internal fragmentation (in the last page)
 - Table space and table maintenance overhead

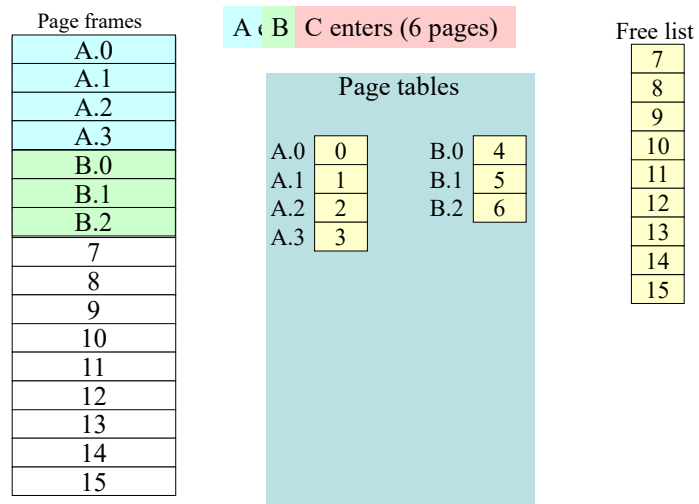
Simple Paging



Simple Paging



Simple Paging



Simple Paging

Page frames	A C B C B leaves	Free list
A.0		13
A.1		14
A.2		15
A.3		
B.0		
B.1		
B.2		
C.0		
C.1		
C.2		
C.3		
C.4		
C.5		
13		
14		
15		

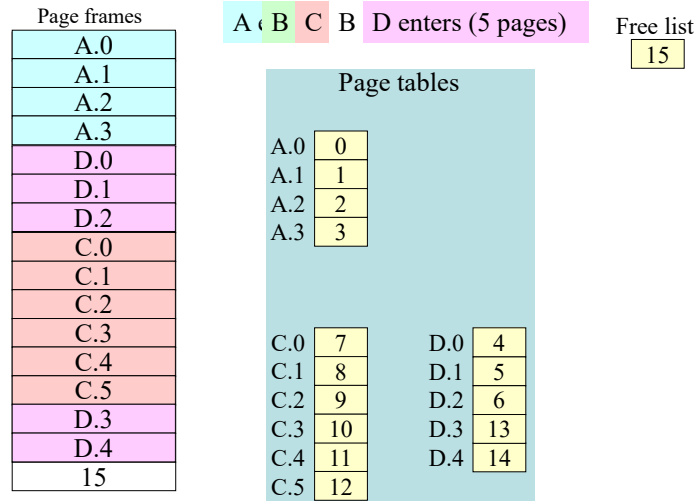
Page tables	
A.0	0
A.1	1
A.2	2
A.3	3
B.0	4
B.1	5
B.2	6
C.0	7
C.1	8
C.2	9
C.3	10
C.4	11
C.5	12

Simple Paging

Page frames	A C B C B D enters (5 pages)	Free list
A.0		4
A.1		5
A.2		6
A.3		13
4		14
5		15
6		
C.0		
C.1		
C.2		
C.3		
C.4		
C.5		
13		
14		
15		

Page tables	
A.0	0
A.1	1
A.2	2
A.3	3
B.0	4
B.1	5
B.2	6
C.0	7
C.1	8
C.2	9
C.3	10
C.4	11
C.5	12

Simple Paging



Addressing in Simple Paging (prelim.)

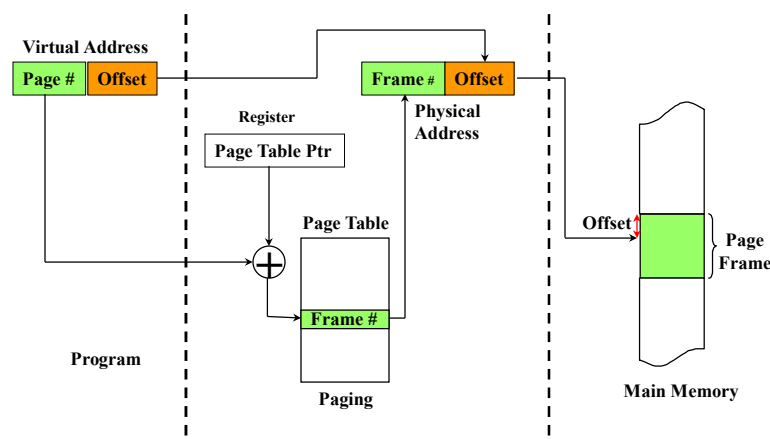
❖ Address computation?

- Logical address x
- Page #: $p = x \text{ div } PS$
 - PS : page size
- Page offset: $t = x \text{ mod } PS$
- Frame # in physical memory: f
 - Need a mapping table, maps p to f
- Memory address = $f * PS + t$
 - $(x \text{ div } PS) * PS + (x \text{ mod } PS)$
- If $PS = 2^n$?
 - n bits for offset and beyond n bits for page number

Addressing in Simple Paging (prelim.)

- ❖ How many bits are required to address the memory
 - 1KB memory = 2^{10} Bytes → 10 bits
 - 1GB memory = 2^{30} Bytes → 30 bits
 - 32 bits → can address 2^{32} Bytes = 4GB
- ❖ Page size and number of pages
 - 1GB memory, 8KB per page ⇒ 128 K pages
- ❖ Address (offset) within a page
 - 8KB per page ⇒ 13 bits to address the offset
- ❖ Address for pages
 - 128 K page ⇒ 17 bits


Addressing in Simple Paging (prelim.)



Addressing in Simple Paging (prelim.)

❖ Addressing

- 1GB memory, 8KB per page \Rightarrow 128K pages



- 00000000000000001010000000101011

- Logical address in hexadecimal: 0x0001402b
- Page number = 1010 = 10
- Page offset = 101011 = 43

P.0	7
P.1	40
P.2	1

- 00000000000000101010000000101011


- Physical address in hexadecimal: 0x0002a02b
- Frame number = 21 = 10101

...	
P.9	15
P.10	21

Addressing in Simple Paging (prelim.)

❖ Addressing

- 64MB memory, 1KB per page \Rightarrow 64K pages



- 0000000000000000000010101101011

- Logical address in hexadecimal: 0x00000aeb
- Page number = 10 = 2
- Page offset = 101101011

P.0	7
P.1	40
P.2	1

- 00000000000000000000101101011


- Physical address in hexadecimal: 0x000006eb
- Frame number = 1

...	
P.9	15
P.10	21

Addressing in Simple Paging (prelim.)

❖ Addressing

- 256MB memory, 4KB per page \Rightarrow 64K pages



- 00000000000000001001010101101011
 - Logical address in hexadecimal: 0x00012aeb
 - Page number = 10010 = 18
 - Exceed page table \rightarrow access violation

P.0	7
P.1	40
P.2	1
	...
P.9	15
P.10	21

Readings

❖ Section 7.1-7.3