Processor Scheduling

What Are We Covering?

- Processor Scheduler
 - > Short Term Schedule ready processes to processor(s)
 - ➤ Long Term Schedule processes to reside in memory
- ❖ How to design a good processor scheduling algorithm
- Performance Metrics

Quality of Scheduling Algorithms

- Performance Metrics
 - > Turnaround Time: average (waiting time + service time)
 - ➤ Response Time
 - Similar to turnaround time
 - Used in interactive environment
 - > Throughput: Number of jobs completed per unit time
 - ➤ Utilization: the percentage of time the processor is busy
 - > Turnaround time / Service time
 - Best matches with users' expectation
 - A longer job can have a longer wait time
 - > Fairness no starvation

Constraints on Scheduling

- Nonpreemptive
 - > Process will run till terminated or blocked
- Preemptive
 - > Process can be interrupted while in execution
 - > E.g., preempted by time quantum
- *Real-Time Scheduling
 - > Each process has to finish before its deadline

Job Characteristics

- CPU bound
 - > Takes long computation time
 - ➤ Little I/O
- ❖ I/O bound
 - ➤ Repeat: run for short time; need I/O

Scheduling Algorithms

- ❖ First-In-First-Out (FIFO, FCFS)
- Shortest Job First
- Highest Response Ratio First
- ❖ Shortest Remaining Time First
- **❖** Round-Robin
- Multilevel Feedback Queue

Scheduling Algorithms

- ❖ First-in-first-out
 - > Easy to implement
 - > Short jobs may be stuck behind a long job
- Shortest Job First
 - ➤ Non-preemptive
 - ➤ Minimize average turnaround time **if** all jobs arrive at the same time
 - A short job still can get stuck behind a long job
 - > Standard deviation of turnaround time may not be good
 - ➤ May cause starvation
 - ➤ Need to predict execution time of each job

Scheduling Algorithms

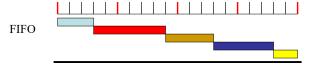
- ❖ Shortest Remaining Time First
 - > Preemptive
 - ➤ Allow a job with shorter remaining time to preempt current running job
 - ➤ Minimize average turnaround time
 - ➤ Similar problems with SJF + jobs need to be preemptable

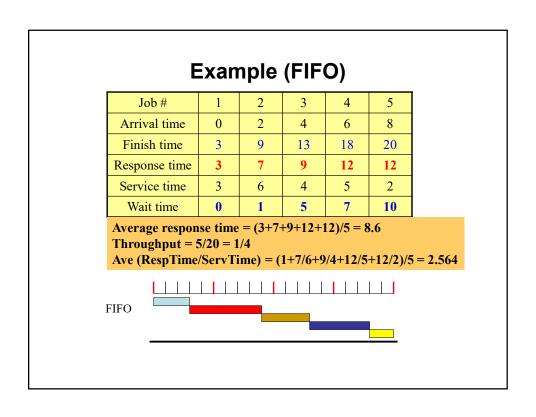
Scheduling Algorithms

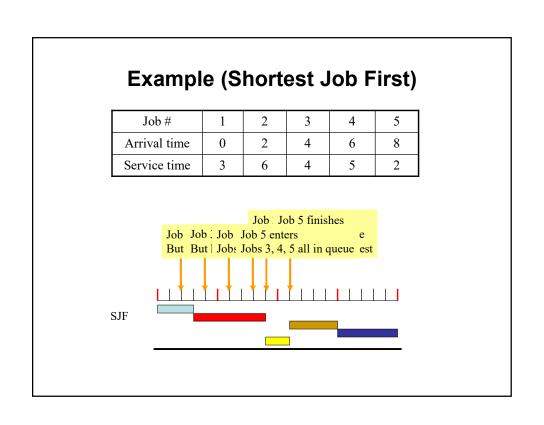
- Highest Response Ratio Next
 - ➤ The ready process with the greatest response ratio (RR) is selected to run
 - > RR = (waiting-time + service-time) / service-time
 - > Shorter jobs are favored
 - Take waiting time in to account
 - Avoid starvation problem
 - More fair
 - ➤ Need to predict execution time of each job
 - Can be non-preemptive or preemptive

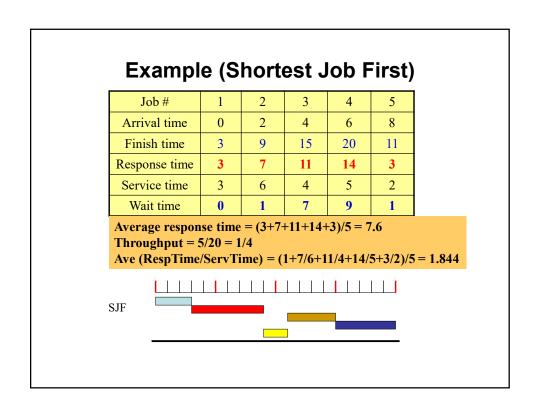
Example (FIFO)

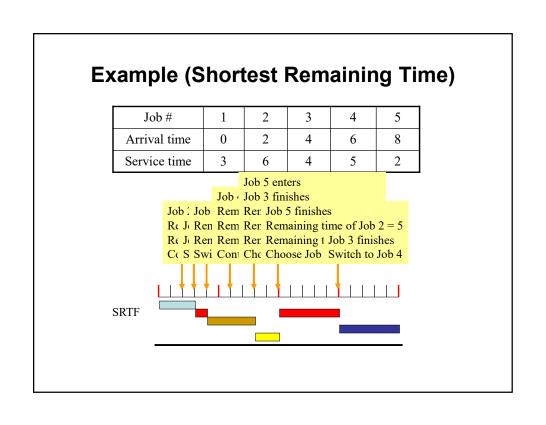
Job#	1	2	3	4	5
Arrival time	0	2	4	6	8
Service time	3	6	4	5	2

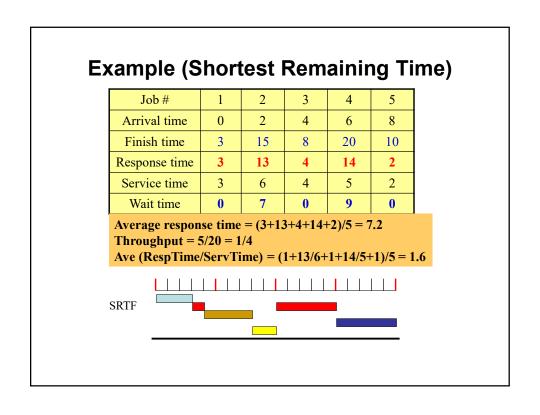


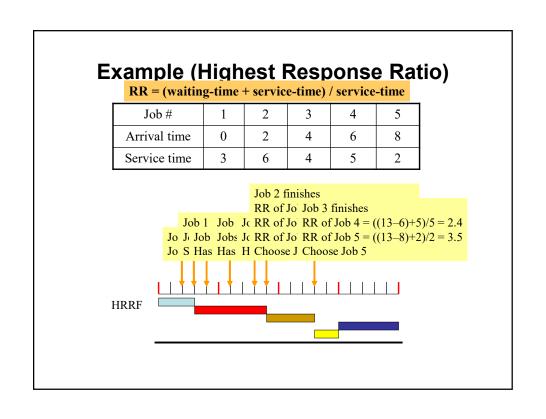












Example (Highest Response Ratio)

Job#	1	2	3	4	5
Arrival time	0	2	4	6	8
Finish time	3	9	13	20	15
Response time	3	7	9	14	7
Service time	3	6	4	5	2
Wait time	0	0	7	9	0

Average response time = (3+7+9+14+7)/5 = 8Throughput = 5/20 = 1/4Ave (RespTime/ServTime) = (1+7/6+9/4+14/5+7/2)/5 = 2.144

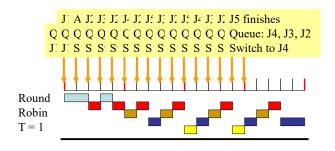
HRRF

Scheduling Algorithms

- *Round-Robin
 - > Process take turns
 - Each executes till its time-quantum expires
 - > Gives relatively reasonable turnaround time
 - Better in standard deviation of turnaround time
 - > Useful when it is hard to estimate the execution time
 - No need to know execution time in advance
 - > Overhead on process switching
 - More process switch ⇒ more context switch
 - ➤ I/O bound jobs may not use up their time quantum
 - Not fair to I/O bound jobs

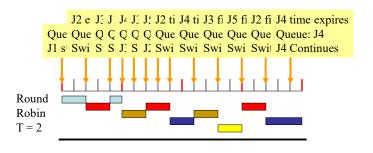
Example (Round Robin, T=1)

Job#	1	2	3	4	5
Arrival time	0	2	4	6	8
Service time	3	6	4	5	2



Example (Round Robin, T=2)

Job#	1	2	3	4	5
Arrival time	0	2	4	6	8
Service time	3	6	4	5	2

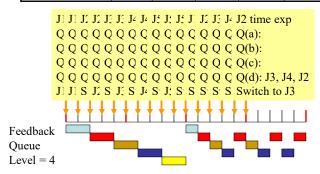


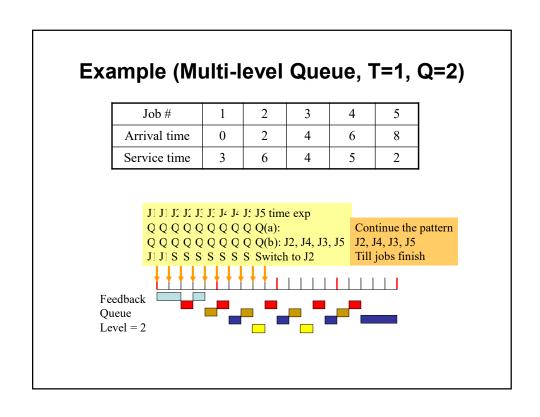
Scheduling Algorithms

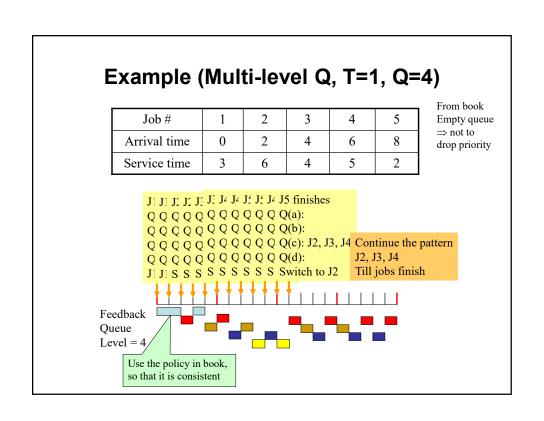
- ❖ Multilevel Feedback Queue (MLFQ)
 - ➤ The system maintains multiple ready queues
 - For example, N priority levels
 - > CPU always serves the jobs in the highest priority queue
 - > A new process will be assigned a high priority
 - ➤ If a process used up its time quantum
 - priority decremented by one (placed in a lower level queue)
 - Favor I/O bound jobs
 - Not use up time quantum → stay at the same priority level
 - ➤ Good for time sharing systems
- ❖ There are many variations of the algorithm

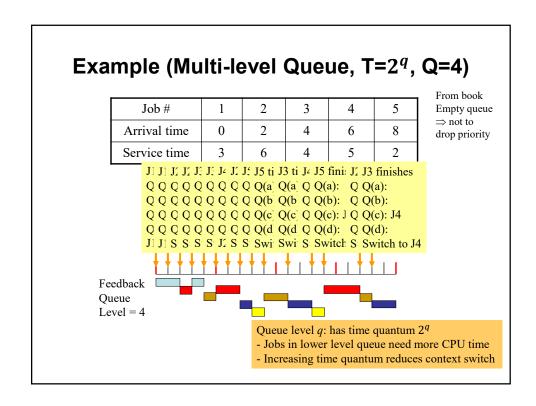
Example (Multi-level Queue, T=1, Q=4)

Job#	1	2	3	4	5
Arrival time	0	2	4	6	8
Service time	3	6	4	5	2









MLFQ-like Scheduling

MLFQ

- ➤ Mostly used algorithm in practice
- ➤ In practice, many variations are used
 - Some consider to move up the task to the next level queue after a certain waiting time *w*
 - A version of Linux considers $q \times T$ as the time quantum for level q queue (T = 10ms)
 - Unix: MLFQ-like, compute priority as many levels of queues
- > Time quantum
 - Old Unix has 1 second time quantum, then changed to 1/10 second, then changed to 10ms

MLFQ-like Scheduling

- MLFQ-like scheduling algorithms
 - > Similar to multilevel feedback queue
 - > Use a number to indicate job priority
 - ➤ Priority level: 0-127
 - ➤ Negative priority value: no interruption
- ❖ A process
 - When enters the system, gets a base priority
 - > Priority changes depending on CPU usage
 - Can also consider waiting time
 - > After time quantum expiration or blocked by IO
 - \Rightarrow Process switch \Rightarrow Recompute priorities of jobs

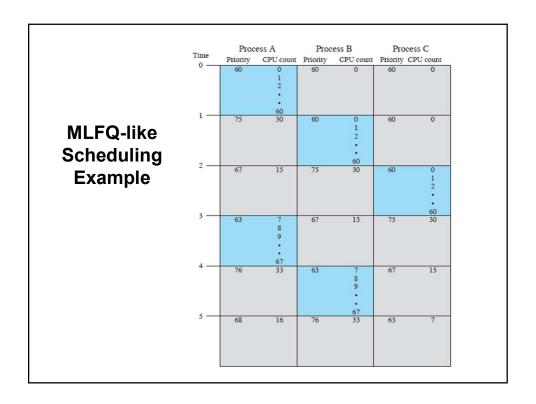
MLFQ-like Scheduling

Priority computation

$$P_{i+1} = base + C_i/2 + nice$$

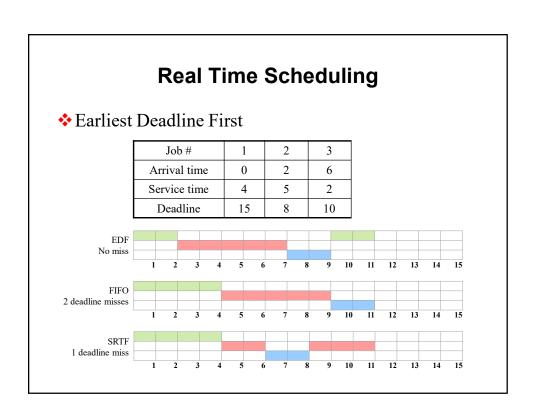
 $C_i = (U_i + C_{i-1})/2$

- \triangleright C_i : a count to keep track of the CPU usage history
- $\triangleright U_i$: CPU time units used in a time quantum i
 - A time quantum is divided into 60 units
- ❖ The priority of a job rises when waiting
 - $\triangleright U_i$ will be 0 when waiting
 - ➤ I/O-bound jobs tend to have higher priorities than CPU-bound jobs

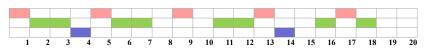


- ❖ Real time systems
 - > Examples of computing that requires real time
 - Monitor refreshing every 16 milliseconds (60Hz): Periodic
 - Traffic light control : Periodical / Aperiodic
 - Image processing in autonomous driving : Aperiodic
 - ➤ Scheduling in RT systems
 - Rate for periodical tasks (how frequent it needs to be activated)
 - Starting time and deadline
 - Preemptive or nonpreemptive
 - Soft or hard deadline

- Fixed Priority
 - Each task has a predetermined priority
 - ➤ Always execute the highest priority task first
 - > Preemptive
- **❖** Earliest Deadline First
 - > The task with the nearest deadline has the highest priority
 - > Preemptive
- *Rate Monotonic Scheduling
 - ➤ Consider periodical tasks
 - Tasks with the highest rates are scheduled first
 - ➤ Preemptive



- *Rate monotonic scheduling
 - > RM schedulability
 - A system is schedulable if $\sum_i u_i \le N(2^{1/N} 1)$
 - $u_i = s_i/p_i$, p_i is the period and s_i is service time
 - *N* is the number of jobs
 - Example: Jobs 1, 2, 3 has (4,1), (5,2), (10,1)
 - $\sum_{i} u_{i} = 1/4 + 2/5 + 1/10 = 0.75 \le 3(2^{\frac{1}{3}} 1) \approx 0.78 \Rightarrow$ Schedulable
 - > Scheduling algorithm
 - Job with shortest period has highest priority ⇒ schedule first



Real Time Scheduling

- Priority based algorithms
 - higher priority task can preempt the execution of lower priority task
 - > Priority inversion
 - A high priority task T_h attempts to preempt a lower priority task T_l \rightarrow but cannot \rightarrow because T_l holds a shared resource that T_h needs
 - ➤ Unbounded priority inversion
 - T_l holds a shared resource that T_h needs, T_l may be preempted by another task T_m (priority of T_h > priority of T_m > priority of T_l) $\Rightarrow T_h$'s wait time is no longer bounded by the resource usage time T_l may incur \Rightarrow In fact, there is no bound \Rightarrow May cause problem
 - NASA rover Mars Pathfinder uses WindRiver VxWorks RTOS, which is susceptible to priority inversion
 - ⇒ Patch the lander software by enabling "**priority inheritance**"

❖ NSAS rover priority inversion

Tasks and events

- T_l: meteorological data gathering task: locked the infor-bus for storing the gathered data
- T_h : infor-bus manager: attempted to lock the infor-bus to read data related to its management task \Rightarrow blocked, waiting for infor-bus
- T_m : communicator, does not use infor-bus, but got scheduled to communicate and preempted T_l 's execution
- ⇒ T_h waits for infor-bus locked by T_l , T_l is ready, waits for T_m for finishing its execution
- T_{other}: watchdog, periodically checks the status of the critical tasks to ensure their proper operation: Noticed that T_h had not been active for some time ⇒ believed that something had gone wrong ⇒ started system reset

Real Time Scheduling

❖ NSAS rover priority inversion

> Problem

- Happened when rover is on Mars ⇒ Took the ground crew many hours to identify the problem and fix it
- In fact, the same problem had occurred during pre-flight tests, but was believed to be a hardware glitch and got put aside

Solution

- Priority inheritance
- If a task T_l blocks the execution of a higher priority task T_h
 ⇒ T_l inherits the priority of T_h (just temporarily)

communicate, does not use into but, but get scheduled to manusciate and promoped T₁ recursion: T₁ was for mile but had tell y T₂ T₃ to ready, wash for T₄, things in constraint manuschilding, periodically shocks the entire of the critical to entire the large operation. Which call that T₁ had not been to it is sense then a young mention. Which call that T₁ had not been to it is sense then a position of the sense that the control of the sense that the sense of the sense of the sense of the transfer of the sense of the sense of the sense of the transfer of the sense of the sense of the transfer of the sense of the transfer of the sense of the transfer of transfer

Readings

- **❖** Sections 9.2, 9.3
 - ➤ Regular CPU scheduling algorithms
- ❖ Section 10.2
 - > Real time scheduling algorithms
- **♦** Sections 10.3-10.6
 - > CPU scheduling in real systems
 - > Optional reading, not covered in class