

Concurrent Programming and Synchronization Problems

What Are We Covering?

- ❖ OS manages resources
 - CPU scheduling
 - Memory management
 - Device management
 - Disk, terminal, clock
- ❖ OS provides file systems
 - File system organization and how to store files on disk
- ❖ **Support concurrent programming**
 - Processes and their states
 - Threads

Concurrent Programming

- ❖ What is concurrent programming?
 - Multiple execution units run concurrently
 - E.g., create new processes
 - E.g., create new threads
 - E.g., run two programs “concurrently”
- ❖ What is the problem?
 - If these units are independent (they don’t interact), then no problem
- ❖ Sometimes, these concurrent units do need to interact

How Programs Interact?

- ❖ Shared memory model (like a blackboard)
 - Multiple threads look at the same information space for their operations

Now: $x = 2$; $y = 3$;
 A: $x = x + 3$;
 B: $y = x - 2$;
 C: $y = y * x$;

$x = 5$; $y = 15$;

What if A and B and C are executed in parallel?

BAC: $x = 5$; $y = 0$;

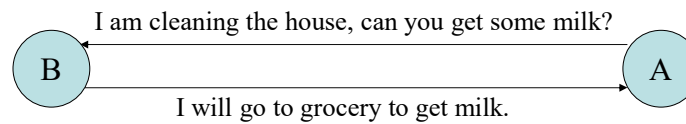
CAB: $x = 5$; $y = 3$;

...

How Programs Interact?

❖ Message passing model

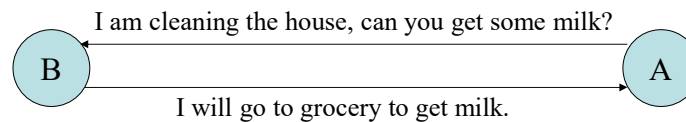
- Multiple processes (threads) have their own space
- They interact by passing messages



What if A and B are not in sync?

How About Synchronization?

❖ Potential Problem in message passing model



A: Sent the message yesterday morning.
 A: Still did not get confirmation in the evening.
 A: Went to the grocery and bought milk.
 B: Receive the message this morning.
 B: Went and bought milk.

Synchronization in Shared Memory

❖ What are the possible outcomes if the two programs run concurrently?

```
compute-A ()
{ .....
  count = count + 2;
  print (count);
}
```

```
compute-B ()
{ .....
  count = count - 3;
  print (count);
}
```

Currently: count = 10

Possible execution orders	Printed results
A1 A2 B1 B2:	12 9
B1 B2 A1 A2:	7 9
A1 B1 A2 B2:	9 9

Any other???

Synchronization in Shared Memory

```
#include <pthread.h>
#include <stdio.h>

int counter;

void *increment (void *arg)
{ int j; int id = *((int *) arg);

  printf ("Thread %d started!\n", id);
  for (j = 0; j < 1000000; j++) counter++;
  printf ("Thread %d is ending!\n", id);
}

int main()
{ int i; pthread_t tid[5];

  for (i = 0; i < 5; i++)
    pthread_create (&tid[i], NULL, increment, (void *) &i);
  for (i = 0; i < 5; i++) pthread_join (tid[i], NULL);

  printf ("Final value of counter = %d\n", counter);
}
```

A bad example of thread programming!!!
How to pass parameters to a thread?
How to protect accesses by multiple threads (counter)?

Synchronization in Shared Memory

❖ Actually, more outputs can occur

```
compute-A ()
{ .....
  count = count + 2;
  print (count);
}
```

⇒

A1 load count
A2 add 2
A3 store count
A4 print (count)

```
compute-B ()
{ .....
  count = count - 3;
  print (count);
}
```

⇒

B1 load count
B2 add -3
B3 store count
B4 print (count)


A1	B1	A2	A3	A4	B2	B3	B4:	12	7
A1	B1	A2	B2	A3	B3	A4	B4:	7	7

Synchronization in Shared Memory

❖ Actually, more outputs can occur

A1 load count
A2 add 2
A3 store count
A4 print (count)

B1 load count
B2 add -3
B3 store count
B4 print (count)



Memory



A exec		load	add				store	print	
B exec				load	add	store			print
AC(A)		10	12				12		
AC(B)				10	7	7			
M	10					7	12		
Screen								12	12

Assumption: print(count): print the memory content of count out (hidden load)

Synchronization in Shared Memory

❖ For the posted thread example

originally count = 0
 loop 10 times
 load count
 add 1
 store count



Context switch in
AC got 5 previously

	loop 5 times	load					add	store	loop 4 times			
			load	add	store	loop 8 times				load	add	store
AC(A)	0 to 5	5					6	6	6-10			
AC(B)			5	6	6	7 to 14				10	11	11
M	0 to 5				6	7 to 14		6	6-10			11

What to Expect?

❖ Which way we want the two threads to execute?

- A first, then B?
- B first, then A?
- One at a time, AB or BA both are OK
- Any execution outcome is fine

❖ How to achieve the desired execution order?

- First case: Initialization: Bwait = 1;
 A: execute A code; Bwait = 0;
 B: while (Bwait) do nothing; execute B code;
- Third case: use a lock box, whoever goes in, lock it
- Last case: nothing needed

What to Expect?

- ❖ Which way we want the two program to execute?
 - One at a time, AB or BA, both are OK
 - This is called **Mutual Exclusion**
- ❖ How to achieve mutual exclusion?
 - Use a lock box, whoever goes in, lock it
 - Within the lock box is a **Critical Region**

What to Expect?

- ❖ Formalizing the concept
 - **Mutual Exclusion:** A requirement for a certain shared object (e.g. code/resources), when satisfied, the system guarantees that no two processes (or other types of agents) will access the object at the same time.
- ❖ How to achieve the desired execution order?
 - **Critical Region:** A section of code, or collection of operations in which only one process may be executing at a given time.
- ❖ Identify critical regions in a program
- ❖ Use **lock, semaphore, or monitor** to achieve mutual execution for the critical regions

What to Expect?

- ❖ Which way we want the two program to execute?
 - A first, then B
 - This is called **synchronization**
- ❖ How to achieve synchronization
 - Use **lock** or **semaphore** to achieve synchronization

Readings

- ❖ Section 5.1