


Deadlock


Deadlock

- ❖ If a group of processes are blocked on their resource requests that can never be satisfied, then it is a deadlock



Dining Philosophers Problem

❖ 5 philosophers, each needs to eat with 2 forks



Philosopher i :

```
repeat
  think;
   $f1 = i$ ;
   $f2 = (i + 1) \bmod 5$ ;
  wait (fork[ $f1$ ]);
  wait (fork[ $f2$ ]);
  eat;
  signal (fork[ $f1$ ]);
  signal (fork[ $f2$ ]);
until false;
```

Necessary Conditions for Deadlock

❖ Mutual Exclusion

- Shared resources are acquired and used mutually exclusively

❖ Hold and Wait

- Each process continues to hold resources already allocated to it while waiting to acquire new resources

❖ No Preemption

- Resources granted to a process can only be released back to the system after the process finished using it


❖ Circular Waiting

- Processes forms a chain in the wait for graph

Wait-For Graph

- ❖ Specify the relationship among processes and resources

- A directed edge from a process to a resource represents that the process is requesting for the resource
- A directed edge from a resource to a process represents that the resource is held by the process



Deadlock Solutions

- ❖ Deadlock Prevention

- Disallow at least one of the four necessary conditions

- ❖ Deadlock Avoidance

- Grant only those requests for available resources which cannot possibly result in deadlock

- ❖ Deadlock Detection

- Grant available resources to requesting processes freely and periodically check for deadlocks
- After deadlock is detected, proceed to **deadlock recovery**

Deadlock Prevention

- ❖ Disallow at least one of the four necessary conditions
- ❖ Avoid Mutual Exclusion
 - Nothing can be done if the resource cannot be shared
- ❖ Allow Preemption
 - Not all resources can be preempted (e.g., printer)
- ❖ Avoid Circular Wait
 - **Linear ordering** on resources
 - E.g., resources A, B, C, $\text{order}(A) = 1$, $\text{order}(B) = 2$, $\text{order}(C) = 3$
 - request (C and A) \Rightarrow request (A); request (C);
 - Problems:
 - Require resource requirement information in advance
 - Low resource utilization
 - Degradation of concurrency (resources are allocated before it is needed)

Deadlock Prevention

❖ Dining philosopher problem

Philosopher i :

```
repeat
  think;
   $f1 = i$ ;  $f2 = (i + 1) \bmod 5$ ;
  if ( $f1 > f2$ ) switch( $f1, f2$ );
  wait (fork[ $f1$ ]);
  wait (fork[ $f2$ ]);
  eat;
  signal (fork[ $f1$ ]);
  signal (fork[ $f2$ ]);
until false;
```

Deadlock Prevention

- ❖ Disallow at least one of the four necessary conditions
- ❖ Avoid Hold-and-Wait
 - Solution 1: A process has to get all of its resources before it begins execution, but the resources may be released at any time, any order
 - Problems for solution 1
 - Require resource requirement information in advance
 - Low resource utilization
 - Degradation of concurrency (worse than linear ordering)
 - Solution 2: A process that has to wait for some resources with: order of one of the held resources > order of the waiting resource
 - Will release all the resources allocated to it
 - (similar to linear ordering without resource information in advance)
 - Problems for solution 2
 - Not all resources can be preempted

Deadlock Avoidance

- ❖ System grants allocations of resources only if it is guaranteed to be deadlock free
 - For a new request requesting for resources
 - System “pretends” to grant the resources
 - Check whether the “new state” is safe
 - Grant the requested resources only if it is safe
 - Safe
 - In the future, the system can allocate resources in at least one order that guarantees free of deadlock
 - Unsafe
 - System may (still may not) end up in deadlock situation
 - Problem
 - Need to know max resource needed for all processes in advance

Banker's Algorithm

❖ System Model

- N : Number of processes
- M : Number of different types of resources
- Input: $Request[i, j]$
 - $1 \leq i \leq N, 1 \leq j \leq M$
 - Number of type j resources being requested by process i
- Output: grant or deny $Request[i, j]$

Banker's Algorithm

❖ Information to be maintained

- $Total[M]$
 - $Total[j]$: Total number of available resources of type j
- $MaxReq[N, M]$
 - $MaxReq[i, j]$: Maximum number of type j resources that may be requested by process i
- $Allocated[N, M]$
 - $Allocated[i, j]$: Number of type j resources that are currently allocated to process i

Banker's Algorithm

❖ Information to be derived

➤ $Need[N, M]$

- $Need[i, j]$: Number of type j resources that may still be needed by process i

➤ $Available[M]$

- $Available[j]$: Number of resources of type j that is still available

Banker's Algorithm Example

❖ Current system state

	Total	Max			Alloc			Need			Avail
		P_1	P_2	P_3	P_1	P_2	P_3	P_1	P_2	P_3	
R_1	7	2	7	3	0	2	1				
R_2	6	3	3	4	1	1	0				
R_3	6	5	1	1	0	1	0				

❖ Request

	P_1	P_2	P_3
R_1	0	1	0
R_2	1	0	0
R_3	0	0	0

Banker's Algorithm Example

❖ Current system state

	Total	Max			Alloc			Need			Avail
		P ₁	P ₂	P ₃	P ₁	P ₂	P ₃	P ₁	P ₂	P ₃	
R ₁	7	2	7	3	0	2	1	2	5	2	4
R ₂	6	3	3	4	1	1	0	2	2	4	4
R ₃	6	5	1	1	0	1	0	5	0	1	5

❖ Request

	P ₁	P ₂	P ₃
R ₁	0	1	0
R ₂	1	0	0
R ₃	0	0	0

Banker's Algorithm

❖ Preliminary checks

- Check to see if the request is invalid
if $Request[i,j] > Need[i,j]$, for any i, j , **then**
error;
- Check to see whether the system has enough resources for the request
if $\sum_i Request[i,j] > Available[j]$, for any j , **then**
deny the request;

Banker's Algorithm

❖ Pretend to allocate

for all j : $Available[j] := Available[j] - \sum_i Request[i,j]$;
for all i, j : $Allocated[i,j] := Allocated[i,j] + Request[i,j]$;
for all i, j : $Need[i,j] := Need[i,j] - Request[i,j]$;

❖ Check safety

if current system is safe **then**
 grant the allocation to the request;
else deny the request

❖ Restore original state if request is denied

for all j : $Available[j] := Available[j] + \sum_i Request[i,j]$;
for all i, j : $Allocated[i,j] := Allocated[i,j] - Request[i,j]$;
for all i, j : $Need[i,j] := Need[i,j] + Request[i,j]$;

Banker's Algorithm

❖ Check safety

for all j : $Temp[j] := Available[j]$; -- not to change Available
 $Pset :=$ the set of all processes;

repeat

if there exists an i such that $Need[i,j] \leq temp[j]$ **then**
 { -- process i can complete its execution
 remove process i from $Pset$;
 for all j : $Temp[j] := Temp[j] + Allocated[i,j]$;
 -- pretend process i is done, return all resources
 }

else return (*unsafe*);

until $Pset = Empty$;

return (*safe*);

Banker's Algorithm Example 1

❖ Current system state

	Total	Max			Alloc			Need			Avl	Request			
		P ₁	P ₂	P ₃	P ₁	P ₂	P ₃	P ₁	P ₂	P ₃		P ₁	P ₂	P ₃	
R ₁	7	2	7	3	0	2	1	2	5	2	4	R ₁	0	1	0
R ₂	6	3	3	4	1	1	0	2	2	4	4	R ₂	1	0	0
R ₃	6	5	1	1	0	1	0	5	0	1	5	R ₃	0	0	0

⇒ no problem

❖ Preliminary checks

➤ if $Request[i,j] > Need[i,j]$, for any i, j
 then **error**;

Request
 0 1 0
 1 0 0
 0 0 2

⇒ error

➤ if $\sum_i Request[i,j] > Available[j]$, for any j
 then **deny**;

Request
 0 5 0
 1 0 0
 0 0 0

⇒ deny

Banker's Algorithm Example 1

❖ Current system state

Total	Max	Alloc	Need	Avl	Request
7	2 7 3	0 2 1	2 5 2	4	0 1 0
6	3 3 4	1 1 0	2 2 4	4	1 0 0
6	5 1 1	0 1 0	5 0 1	5	0 0 0

❖ After “pretend to allocate”

Total	Max	Alloc	Need	Avl
7	2 7 3	0 3 1	2 4 2	3
6	3 3 4	2 1 0	1 2 4	3
6	5 1 1	0 1 0	5 0 1	5

No change

Banker's Algo Example 1 (Check Safety)

❖ Check safety

Total	Max	Alloc	Need	Avl	
7	2 7 3	0 3 1	2 4 2	3	Which process can continue?
6	3 3 4	2 1 0	1 2 4	3	$Need[i,j] \leq Avail[j]$
6	5 1 1	0 1 0	5 0 1	5	P_1 can!
<hr/>					Pretend:
	7 3	3 1	4 2	3	- Allocate resources to P_1
	3 4	1 0	2 4	5	- P_1 finishes
	1 1	1 0	0 1	5	- P_1 returns all its resources
<hr/>					
	7	3	4	4	P_3 can continue!
	3	1	2	5	P_2 can continue!
	1	1	0	5	All finishes \Rightarrow Safe!
					\Rightarrow Grant request!

Banker's Algorithm Example 2

❖ Current system state

Total	Max	Alloc	Need	Avl	Request
7	2 7 3	0 2 1	2 5 2	4	0 1 0
6	3 3 4	1 1 0	2 2 4	4	1 0 0
6	5 1 1	0 1 0	5 0 1	5	0 0 1

❖ After "pretend to allocate"

Total	Max	Alloc	Need	Avl	
7	2 7 3	0 3 1	2 4 2	3	Which process can continue?
6	3 3 4	2 1 0	1 2 4	3	$Need[i,j] \leq Avail[j]$
6	5 1 1	0 1 1	5 0 0	4	None of them can!!!!
					\Rightarrow Deny request

Banker's Algorithm Example 3

❖ Current system state

Total	Max	Alloc	Need	Avl	Request
5	1 4 3	0 1 0	1 3 3	4	0 1 0
4	4 3 1	0 2 0	4 1 1	2	0 1 1
3	2 1 1	2 0 0	0 1 1	1	0 0 0

❖ After "pretend to allocate"

Total	Max	Alloc	Need	Avl
5	1 4 3	0 2 0	1 2 3	3
4	4 3 1	0 3 1	4 0 0	0
3	2 1 1	2 0 0	0 1 1	1

Banker's Algorithm Example 3

❖ Check safety

Total	Max	Alloc	Need	Avl	Which process can continue?
5	1 4 3	0 2 0	1 2 3	3	$Need[i,j] \leq Avail[j]$
4	4 3 1	0 3 1	4 0 0	0	P₂ and P₃ both can!
3	2 1 1	2 0 0	0 1 1	1	Choose P₂
<hr/>					Pretend:
5		0 0	1 3	5	- Allocate resources to P ₂
4		0 1	4 0	3	- P ₂ finishes
3		2 0	0 1	1	- P ₂ returns all its resources
<hr/>					P ₃ can continue!
5		0	1	5	P ₃ returns its resources!
4		0	4	4	All finishes ⇒ Safe!
3		2	0	1	⇒ Grant request!

Banker's Algorithm Example 4

❖ Current system state

Total	Max	Alloc	Need	Avl	Request
5	1 4 3	0 1 0	1 3 3	4	1 0 0
4	4 2 4	0 1 1	4 1 3	2	0 0 0
3	2 1 2	1 0 0	1 1 2	2	1 0 0

❖ After “pretend to allocate”

Total	Max	Alloc	Need	Avl	Which process can continue? <i>Need[i, j] ≤ Avail[j]</i>
5	1 4 3	1 1 0	0 3 3	3	
4	4 2 4	0 1 1	4 1 3	2	P₂ can ⇒ return all
3	2 1 2	2 0 0	0 1 2	1	None of the rest can ⇒ Deny request

Deadlock Detection

❖ Just let deadlock happen

- No prevention, no avoidance

❖ System periodically checks for deadlock

- When a deadlock is detected, a recovery process is required to break the cycle
- Use the safety check algorithm for deadlock detection
 - No need to know maximum resource requirements “Max”
 - No need to know still needed resources “Needed”

❖ deadlock recovery

- Kill one process at a time till deadlock cycle is broken

Safety Algorithm for Detection

❖ Check safety

	Total	Alloc			Avl	Req			
		P ₁	P ₂	P ₃		P ₁	P ₂	P ₃	
R ₁	3	0	2	0	1	2	0	1	P ₂ can get the resources → P ₂ return resources
R ₂	4	0	1	1	2	0	2	1	
R ₃	2	1	0	0	1	0	1	2	
<hr/>									
R ₁	3	0		0	3	2		1	P ₁ can get the resources → P ₁ return resources
R ₂	4	0		1	3	0		1	
R ₃	2	1		0	1	0		2	
<hr/>									
R ₁	3			0	3			1	P ₃ can get the resources All finishes ⇒ Safe! ⇒ No deadlock!
R ₂	4			1	3			1	
R ₃	2			0	2			2	

Safety Algorithm for Detection

❖ Check safety

	Total	Alloc			Avl	Req			
		P ₁	P ₂	P ₃		P ₁	P ₂	P ₃	
R ₁	2	0	0	1	1	2	0	1	P ₂ can get the resources → P ₂ return resources
R ₂	4	0	1	1	2	0	2	1	
R ₃	2	1	0	0	1	0	1	2	

R ₁	2	0	1	1	1	2	1		No process can get the resources ⇒ Deadlock P ₁ and P ₃ involved
R ₂	4	0	1	3	0	1			
R ₃	2	1	0	1	0	2			

Comparisons

- ❖ Overhead in deadlock prevention (linear ordering)
 - Take resources in advance \Rightarrow Lower resources utilization; Reduce degree of concurrency
- ❖ Overhead in deadlock avoidance
 - Need to update resource information for every allocation/deallocation
 - Need to run safety algorithm for each request (very expensive)
- ❖ Overhead in deadlock detection
 - Detection algorithm can be expensive
 - but only when there is a potential deadlock
 - Need to wait for resources till a deadlock is detected
- ❖ No perfect solution -- Consider combinations

Readings

- ❖ Section 6.2-6.4, 6.6
- ❖ Good but not covered: 6.8