

Synchronization Methods in Shared Memory Model

Mutual Exclusion and Synchronization

❖ Mutual exclusion

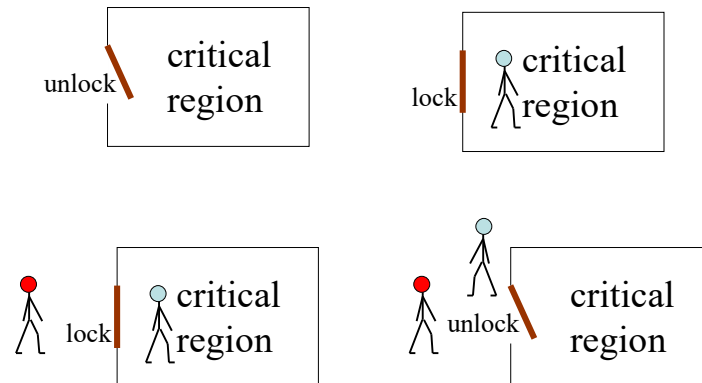
- Guarantee that no two processes (or other types of agents) should access a critical region at the same time

❖ Synchronization

- Guarantee that one step should occur before/after another
- E.g., P_1 and P_2 run concurrently
 - P_1 : S_{11}, S_{12}
 - P_2 : S_{21}, S_{22}
 - Guarantee that S_{12} executes before S_{22}

❖ Use **lock, semaphore, or monitor** to achieve mutual execution and/or synchronization

Critical Region and Locking



Simply set the lock -- Test-Set

❖ Consider two processes: A and B

Process A and B: *lock* is initialized to *false*

while *lock* **do** nothing;

lock := *true*;

< critical section >

lock := *false*;

➤ Execution

A: **while** *lock* **do** nothing; -- find *lock* = *false*

B: **while** *lock* **do** nothing; -- find *lock* = *false*

A: *lock* := *true*;

B: *lock* := *true*;

A: < critical section >

B: < critical section >

➤ does not guarantee mutual exclusion at all

Too bad!!!

Simply set the lock -- Set-Test

❖ Consider two processes: A and B

Process A and B:

```
lock := true;
while lock do nothing;
< critical section >
lock := false;
```

➤ Execution

```
A: lock := true;           B: lock := true;
A: while lock do nothing; -- find lock = true
B: while lock do nothing; -- find lock = true
```

➤ Deadlock, neither can enter CS

Still bad!!!

Test-Set, Use separate locks

Process A:

```
while lock[B] do nothing;
lock[A] := true;
< critical section >
lock[A] := false;
```

Process B:

```
while lock[A] do nothing;
lock[B] := true;
< critical section >
lock[B] := false;
```

Initialization: $lock[A] = lock[B] = false$

➤ Execution

```
A: while lock[B] do nothing; -- find lock[B] = false
B: while lock[A] do nothing; -- find lock[A] = false
A: lock[A] := true;
B: lock[B] := true;
A: < critical section >   B: < critical section >
```

➤ Similar to single lock algorithm, does not guarantee mutual exclusion

Set-Test, Use separate locks

Process A:

lock[B] = F ⇒
no competition
⇒ Enter CS

```
lock[A] := true;
while lock[B] do nothing;
< critical section >
lock[A] := false;
```

Process B:

```
lock[B] := true;
while lock[A] do nothing;
< critical section >
lock[B] := false;
```

➤ Execution

A: *lock[A] := true;*

B: *lock[B] := true;*

A: **while** *lock[B]* **do** nothing; -- find *lock[B] = true*

B: **while** *lock[A]* **do** nothing; -- find *lock[A] = true*

➤ Similar to the single lock case, deadlock, neither can enter

➤ But, if there is no completion, no problem

Improved from the single lock case

Set-Test, Separate lock, Give away

Process A:

Give away on
competition

```
lock[A] := true;
while lock[B] do nothing
{ lock[A] := false;
  delay (short time);
  lock[A] := true; }
< critical section >
lock[A] := false;
```

Process B:

```
lock[B] := true;
while lock[A] do {
  lock[B] := false;
  delay (short time);
  lock[B] := true; }
< critical section >
lock[B] := false;
```

➤ Execution

A: *lock[A] := true;*

B: *lock[B] := true;*

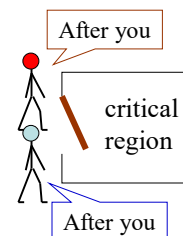
A: **while** *lock[B]* **do** -- find *lock[B] = true*

B: **while** *lock[A]* **do** -- find *lock[A] = true*

A: *lock[A] := false;*

B: *lock[B] := false;*

➤ Close, but has mutual courtesy problem



Take turns

❖ Instead of using two locks, use one turn variable

Process A:

```
while turn != A do nothing;
< critical section >
turn := B;
```

Process B:

```
while turn != B do nothing;
< critical section >
turn := A;
```

- Guarantees mutual exclusion
- But processes have to take turns to enter CS
 - If one process does not have its turn (failed or not needing it), the others will not get it
 - More processes → taking turn is a bigger problem
- Close to a solution

Set-Test, Separate lock + Take turns

Original code for
Set-Test, separate lock, give away

```
Process A:
  lock[A] := true;

  while lock[B] do
  {
    lock[A] := false;
    delay (short time);
    lock[A] := true;
  }
  < critical section >

  lock[A] := false;
```

if *lock*[*B*] = *false*
⇒ no competition
⇒ Enter CS

When compete
⇒ Take turns

Original code for **Take turns**

```
Process A:
  while turn != A do nothing;
  < critical section >
  turn := B;
```

Set-Test, Separate lock + Take turns

Original code for **Take turns**
Set-Test, separate lock, give-away

Process A:
`lock[A] := true;`
while `lock[B]` **do**
 check turn
 { `lock[A] := false;`
 delay (short time);
 `lock[A] := true;`
 }
 < critical section >
turn := B;
`lock[A] := false;`

if `lock[B] = false`
 ⇒ no competition
 ⇒ Enter CS
 When compete
 ⇒ Take turns

Original code for **Take turns**

Process A:
while `turn != A` **do** nothing;
 < critical section >
`turn := B;`

if `turn = A` (`lock[A]` has been set)
 (B cannot proceed) ⇒ A goes into CS
if `turn != A`, but **!lock[B]**
 (B does not compete) ⇒ A goes into CS
else loop till `lock[B] = false` or `turn = A`
 (after B finishes, this will be true)

Set-Test, Separate lock + Take turns

Original code for **Take turns**
Set-Test, separate lock, give-away

Process A:
`lock[A] := true;`
while `lock[B]`
 and `turn != A` **do**
 { ~~`lock[A] := false;`~~
 delay (short time);
 ~~`lock[A] := true;`~~
 }
 < critical section >
turn := B;
`lock[A] := false;`

No need to
 release `lock[A]`
 When `turn = B`, B
 proceeds anyway

Original code for **Take turns**

Process A:
while `turn != A` **do** nothing;
 < critical section >
`turn := B;`

if `turn = A`
 (B cannot proceed) ⇒ A goes into CS
if `turn != A`, but **!lock[B]**
 (B does not compete) ⇒ A goes into CS
else loop till `lock[B] = false` or `turn = A`
 (after B finishes, this will be true)

Set-Test, Separate lock + Take turns

Original code for **Take turns**
Set-Test, separate lock, give-away

```

Process A:
  lock[A] := true;
  turn := B;
  while lock[B]
  and turn != A do
  { lock[A] := false;
    delay (short time);
    lock[A] := true;
  }
  < critical section >
  turn := B;
  lock[A] := false;

```

Move **turn := B** up
When B comes to compete CS
⇒ B sets **turn := A**
⇒ Problem solved!!!

if turn = A
(B cannot proceed) ⇒ A goes into CS
if turn != A, but !lock[B]
(B does not compete) ⇒ A goes into CS
but B may set lock[B] and turn was B
else loop till lock[B] = false or turn = A
(after B finishes, this will be true)

Both enter CS

Set-Test, Separate lock + Take turns

Original code for **Take turns**
Set-Test, separate lock, give-away

```

Process A:
  lock[A] := true;
  turn := B;
  while lock[B]
  and turn != A do
  { lock[A] := false;
    delay (short time);
    lock[A] := true;
  }
  < critical section >
  turn := B;
  lock[A] := false;

```

Move **turn := B** up
When B comes to compete CS
⇒ B sets **turn := A**
⇒ Problem solved!!!

Another benefit from the move
Both yield the turn
Later comer succeeds in yield
Earlier comer wins the turn

Simplification of Lock and Take Turns

Process A:

```
lock[A] := true;
turn := B;
while lock[B] and
  (turn != A) do nothing;
< critical section >
lock[A] := false;
```

Process B:

```
lock[B] := true;
turn := A;
while lock[A] and
  (turn != B) do nothing;
< critical section >
lock[B] := false;
```

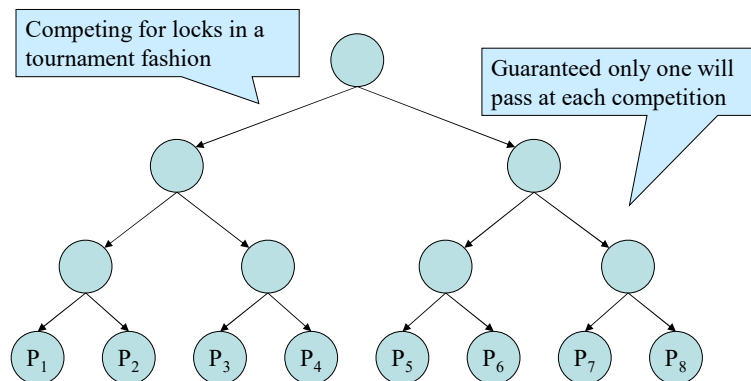
This is the Peterson's algorithm

Before Peterson's algorithm, there was a similar but more complex solution, the **Dekker's algorithm**. But it is easier to understand this one by stepping to it directly.

❖ Software solutions

- Quite significant overhead for a simple lock

Extend Lock for N Processes



Hardware Solutions

❖ Test and Set Instruction

- In one instruction cycle: $Reg \rightarrow lock$ and $lock \rightarrow Reg$
- = swap (register, memory location)

❖ Use *Test-and-Set* for mutual exclusion

$R := true;$

tset instruction in x86

repeat *Test-and-Set* ($R, lock$);

until $R = false;$

< critical section >;

$lock := false;$

- Similar to the first algorithm, but perform check and set in one instruction cycle

Use of Lock

❖ Examples

- How to use *lock* for mutual exclusion?

$lock(lck);$

< critical section >

$unlock(lck);$

- How to use *lock* for synchronization?

▪ E.g.: $P_1: S_{11}, S_{12}$ $P_2: S_{21}, S_{22}$

– Guarantee that S_{12} executes before S_{22}

$P_1: S_{11}; S_{12}; unlock(lck);$

$P_2: S_{21}; lock(lck); S_{22}; unlock(lck);$

- What should be the initial value of *lck* (locked/unlocked)

❖ E.g.: A, B executed mutual exclusively, D after A, B

Use of Lock

❖ How to use *lock* for mutual exclusion?

➤ Example: requirement

- A and B executed mutual exclusively
- D after A and B

➤ Solution

- `lock(mutex); A; unlock(mutex); unlock(sync-a);`
- `lock(mutex); B; unlock(mutex); unlock(sync-b);`
- `lock(sync-a); lock(sync-b); D`
 - How should the locks be initialized?
 - Can we use one sync lock?

Problem with Lock

❖ Software solutions are slow

❖ Hardware solutions

- Much more efficient
- Most of the systems support hardware solutions
 - Test-and-Set, Disable interrupts

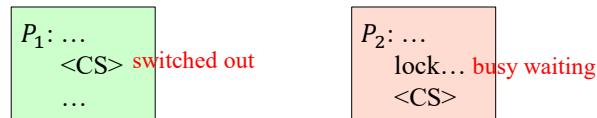
❖ Lock problems in general

- Require busy waiting
 - Also called spin lock
- Unfair
 - First comer may not get the lock first
 - Potential starvation problem

Problem with Lock

❖ Busy waiting problem

- Consider that P_1 is in critical section and P_2 waits for it
- Single core systems
 - P_1 got switched out before finishing the critical section
 - P_2 got switched in, execute *lock*, start to wait on it
 - During P_2 's time quantum, P_1 cannot execute and release *lock*
 $\Rightarrow P_2$ will busy-wait on *lock* till its time quantum expires
- Multiple core systems
 - Same problem: P_1 can still get switched out no matter which processor it was running on and P_2 still will busy-wait on *lock*



Disable Interrupt

❖ How to solve the problem with lock

- Disable Interrupts (another hardware solution)
 - Disable interrupts before CS and enable them after CS
- Disable interrupts;
 < critical section >
 Enable interrupts;
- No process switch when executing CS
- How about in multi-core systems?

IF in x86: interrupt flag = 1 handle interrupts = 0 disable interrupts CLI instruction: clear IF SLI instruction: set IF	But IF is only for hardware interrupts CLI & SLI are privileged instructions (kernel only)
---	---

Semaphores

❖ How to solve the problems with locks

- Use a queue to block processes waiting to enter CS
 - Avoid busy waiting: processes are blocked until its turn comes
 - Achieve fairness: queue enforces order

❖ Integer semaphore

- Has a queue: to hold blocked processes
- Has a counter
 - Control: > 0 : pass; ≤ 0 : blocked
 - Counting the number of processes in queue: ≤ 0
- Offer Two functions: *wait* and *signal*
 - *wait*: request to enter the critical section - also P
 - *signal*: release the critical section - also V

Semaphore Implementation

❖ **type semaphore =**

```
record count: integer;
        queue: queue of processes;
```

```
end;
```

```
var s: semaphore := 10;
```

❖ *wait(s):*

```
s.count := s.count - 1;
if (s.count < 0) then
  { block process P;
    put P in s.queue;
  }
```

signal(s):

```
s.count := s.count + 1;
if (s.count <= 0) then
  { remove P from s.queue;
    put P in ready list;
  }
```

Semaphore Implementation

```

❖ type semaphore
    record
        count: integer;
    end;
var s: semaphore;

❖ wait(s):
    s.count := s.count - 1;
    if (s.count < 0) then
        { block process P;
          put P in s.queue;
        }

    signal(s):
        s.count := s.count + 1;
        if (s.count <= 0) then
            { remove P from s.queue;
              put P in ready list;
            }

```

A: load s.count -- s.count = 1
 B: load s.count -- s.count = 1
 A: add -1
 B: add -1
 A: (s.count < 0) → no → <pass semaphore wait>
 B: (s.count < 0) → no → <pass semaphore wait>
Too bad!!!

Semaphore Implementation

- ❖ Semaphore *wait* and *signal* functions are themselves critical sections
 - Simultaneous accesses to *s.count* and *s.queue* can cause problem
- ❖ Need to use lock to protect *wait* and *signal* functions
 - Lock solution has busy waiting problem
 - Busy waiting for small segments of code is not too much overhead (very low probability of getting switched out)

Semaphore Implementation

<pre>wait(s): lock (s.lck); s.count := s.count - 1; if (s.count < 0) then { block process P; put P in s.queue; } unlock (s.lck);</pre>	<pre>signal(s): lock (s.lck); s.count := s.count + 1; if (s.count <= 0) then { remove P from s.queue; put P in ready list; } unlock (s.lck);</pre>
---	---

Semaphore has
s.count, *s.queue*, and *s.lck*

Use of Semaphores

❖ Use semaphores for mutual exclusion

```
s.count := 1;
wait (s);
< critical section >
signal (s);
```

❖ Use semaphores for synchronization

➤ E.g., Guarantee that S_{12} executes before S_{22}

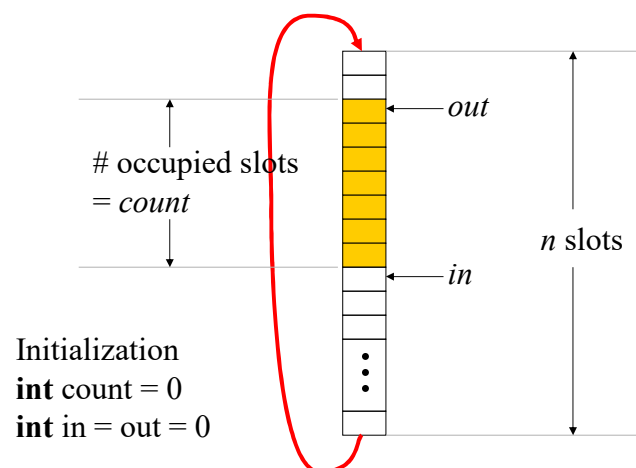
```
P1: S11, S12          P2: S21, S22
s.count := 0;
P1: S11, S12, signal (s);
P2: S21, wait (s); S22
```

Use of Semaphores

❖ Typical semaphore programming examples

- Bounded buffer problem
- Reader-writer problem
- Bakery problem

Bounded Buffer Problem



Bounded Buffer Problem

process producer;

repeat

produce *item*;

while *count* = *n* **do** nothing;

buffer[*in*] := *item*;

in := (*in*+1) mod *n*;

count := *count* + 1;

until forever

process consumer;

repeat

while *count* = 0 **do** nothing;

item := *buffer*[*out*];

out := (*out*+1) mod *n*;

count := *count* - 1;

consume *item*;

until forever

Bounded Buffer Problem

process producer;

repeat

produce *item*;

while *count* = *n* **do** nothing;

buffer[*in*] := *item*;

in := (*in*+1) mod *n*;

count := *count* + 1;

until forever

If there are two producers

A1: *buffer*[5] := *item*_A;

A2: *buffer*[5] := *item*_B;

Put in the same slot!!!

while *count* = 0 **do** nothing;

item := *buffer*[*out*];

out := (*out*+1) mod *n*;

count := *count* - 1;

Current: *count* = 9; *n* = 10

A: load *count* -- *count* = 9

B: load *count* -- *count* = 9

A: incr *count* -- *count* = 10

B: decr *count* -- *count* = 8

Wrong count!!!

Bounded Buffer Problem

<pre> process producer; repeat produce <i>item</i>; <i>wait (mutex);</i> while <i>count</i> = <i>n</i> do nothing; <i>buffer[in] := item;</i> <i>in := (in+1) mod n;</i> <i>count := count + 1;</i> <i>signal (mutex);</i> until forever </pre>	<pre> process consumer; repeat <i>wait (mutex);</i> while <i>count</i> = 0 do nothing; <i>item := buffer[out];</i> <i>out := (out+1) mod n;</i> <i>count := count - 1;</i> <i>signal (mutex);</i> consume <i>item</i>; until forever </pre>
--	--

Bounded Buffer Problem

<pre> process producer; repeat produce <i>item</i>; <i>wait (mutex);</i> while <i>count</i> = <i>n</i> do nothing; <i>buffer[in] := item;</i> <i>in := (in+1) mod n;</i> <i>count := count + 1;</i> <i>signal (mutex);</i> until forever </pre>	<pre> <i>item := buffer[out];</i> <i>out := (out+1) mod n;</i> <i>count := count - 1;</i> <i>signal (mutex);</i> consume <i>item</i>; until forever </pre>
--	---

Current: count = 10; n = 10;
 P: get the lock
 P: find count = 10, wait
 ⇒
 Need to enter the critical section to change count
 But no one can enter CS

Bounded Buffer Problem

<pre> process producer; repeat produce <i>item</i>; while <i>count</i> = <i>n</i> do nothing; <i>wait (mutex);</i> <i>buffer[in] := item;</i> <i>in := (in+1) mod n;</i> <i>count := count + 1;</i> <i>signal (mutex);</i> until forever </pre>	<pre> process consumer; repeat while <i>count</i> = 0 do nothing; <i>wait (mutex);</i> <i>item := buffer[out];</i> <i>out := (out+1) mod n;</i> <i>count := count - 1;</i> <i>signal (mutex);</i> consume <i>item</i>; until forever </pre>
--	--

Bounded Buffer Problem

<pre> process producer; repeat produce <i>item</i>; while <i>count</i> = <i>n</i> do nothing; <i>wait (mutex);</i> <i>buffer[in] := item;</i> <i>in := (in+1) mod n;</i> <i>count := count + 1;</i> <i>signal (mutex);</i> until forever </pre>	<p>Current: <i>count</i> = 9; <i>n</i> = 10</p> <p>P1: check (<i>count</i> = <i>n</i>) -- <i>count</i> = 9, not full</p> <p>P2: check (<i>count</i> = <i>n</i>) -- <i>count</i> = 9, not full</p> <p>P2: get the mutex put item in buffer <i>count</i> = <i>count</i> + 1; -- buffer full</p> <p>P1: checked earlier, buffer was not full, will not check again, but buffer is full, cause problem</p> <p>⇒ Put the while loop in</p> <p>until forever</p>
--	--

Bounded Buffer Problem

❖ Check full or empty

- Producer check “buffer full”
 - Ok to block other producers but should not block consumers
- Consumer check “buffer empty”
 - Ok to block other consumers but should not block producers

❖ How about using different semaphores?

- Update buffer/count: protected by the same semaphore
- Check full: one semaphore to protect from all producers
- Check empty: one semaphore to protect from all consumers

Bounded Buffer Problem

Use semaphores for conditional wait \Rightarrow Solve the problem + avoid busy waiting

process producer;

repeat

produce *item*;

~~**while** *count* = *n* **do nothing;**~~

wait (mutex);

buffer[in] := item;

in := (in+1) mod n;

count := count + 1;

signal (mutex);

until forever

wait on full condition
(no longer empty)

process consumer;

repeat

wait on empty condition

~~**while** *count* = 0 **do nothing;**~~

wait (mutex);

item := buffer[out];

out := (out+1) mod n;

count := count - 1;

signal (mutex);

consume *item*;

until forever

Signal full
(no longer full)

Semaphores need to be signaled. When to signal???

When the condition no long holds!

Bounded Buffer Problem

Need to initialize 3 semaphores!

❖ Use semaphore count to do counting

```

process producer;
  loop
    produce item;
    wait (full);
    wait (mutex);
    access shared_buffer;
    signal (mutex);
    signal (empty);
  end
end;
```

```

process consumer;
  loop
    wait (empty);
    wait (mutex);
    access shared_buffer;
    signal (full);
    signal (mutex);
    consume item;
  end
end;
```

Bounded Buffer Problem

Need to initialize 3 semaphores!

❖ Use semaphore count to do counting

```

process producer;
  loop
    produce item;
    wait (full);
    wait (mutex);
    access shared_buffer;
    signal (mutex);
    signal (empty);
  end
end;
```

init: full = N
-- distance to full

Init:
mutex = 1

```

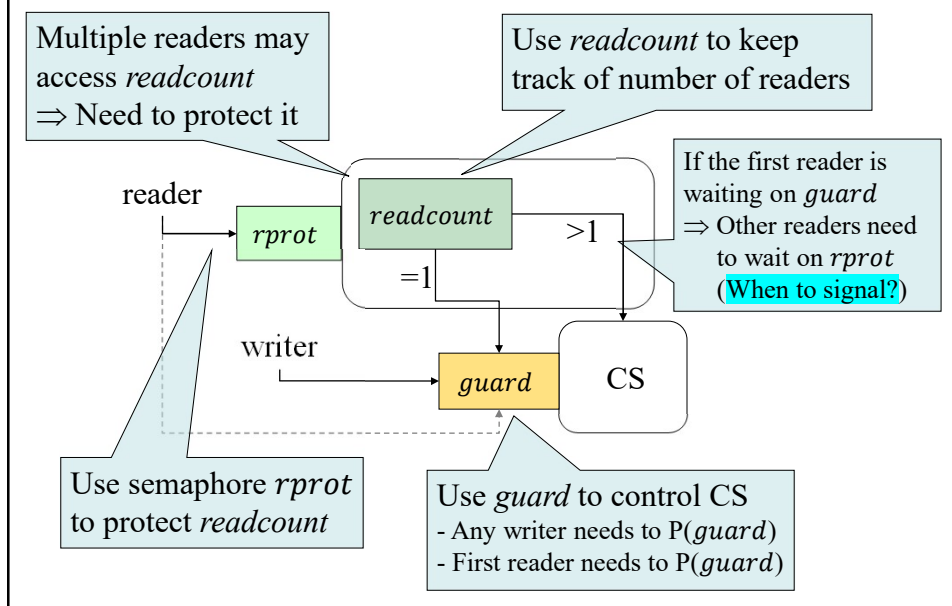
process consumer;
  loop
    wait (empty);
    wait (mutex);
    access shared_buffer;
    signal (full);
    signal (mutex);
    consume item;
  end
end;
```

init: empty = 0
-- distance to empty

Reader-Writer Problem

- ❖ Many readers can read at the same time without causing problem
- ❖ When there is a writer, it has to write exclusively
- ❖ First reader/writer problem (reader has priority)
 - If ≥ 1 reader in CS \Rightarrow allow more readers to get in
 - If a writer is in CS, no one else can enter
 - This way readers has a higher priority
 - Have starvation problem for writers
- ❖ There are other forms of the reader/writer problem

First Reader-Writer Problem



First Reader-Writer Problem

❖ Entry time

- $readcount = 1 \rightarrow$ first reader \Rightarrow wait on *guard*
 - If there is a writer in CS: this reader waits on *guard*
 - After entering CS: signal *rprot*, to let other readers access *readcount* and CS
- $readcount > 1 \rightarrow$ at least one reader is already in CS
 - Proceed to CS directly

❖ Exit time

- $readcount > 0 \rightarrow$ more readers in CS \Rightarrow Just leave
- $readcount = 0 \rightarrow$ no more readers in CS \Rightarrow signal *guard*
 - If there is a writer, signal *guard* to let a writer proceed
 - Signal simply allows future entrance to CS

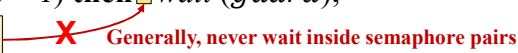
First Reader-Writer Problem

❖ Reader Process:

```

wait (rprot);
  readcount := readcount + 1;
  if (readcount = 1) then wait (guard);
  signal (rprot);
  < enter critical section for read >
wait (rprot);
  readcount := readcount - 1;
  if readcount = 0 then signal (guard);
signal(rprot);

```



First Reader-Writer Problem

❖ Writer Process:

```
wait (guard);  
< enter critical section for write >  
signal (guard);
```

Bakery Problem

- ❖ A bakery has N salesmen
- ❖ Customers can come at any time
 - No bound on number of customers
- ❖ A salesman waits for a customer and provide service
 - wait for customer; provide service;
- ❖ If no salesman is available, a customer has to wait
 - wait for salesman; get service;
- ❖ Write a program using semaphores to synchronize the salesmen and customers
 - cust: customer
 - sales: salesman

Bakery Problem

❖ Consider a different version first

- There are N salesman and M customers

Salesman Process:

```

repeat
  wait (cust);
  provide service;
  signal (sales);
until false;
  
```

Customer Process:

```

wait (sales);
get service;
signal (cust);
  
```

But
Customers are not the same and
are of unknown number

Customer Process:

```

signal (cust);
wait (sales);
get service;
  
```

Initialization
cust = 0
sales = N

Types of Semaphores

❖ Integer semaphore (general semaphore): has a count

- What we used are integer semaphores
- Most common type

❖ Binary semaphores

- Does not keep count, semaphore value can only be 0 or 1
 - Similar to lock, but has a queue
- In some cases, a binary semaphore is sufficient; e.g., mutex

Properties of Semaphores

❖ Lock

- Has the busy waiting problem

❖ Semaphore

- Requires context switch

❖ When the critical section is small

- Short busy-wait may be cheaper than context switch
⇒ use lock instead of semaphore
- But remember the potential of: process in CS is switched out and other processes waiting on the lock will busy wait till time quantum expires
 - Hopefully the probability of this is very small
 - And, semaphore contains lock and has the same problem

Semaphores in Pthread

❖ Creating a semaphore

```
#include <semaphore.h>
sem_t sem;
```

❖ Operations on a semaphore

- `int sem_init (sem_t *sem, int pshared, unsigned int value);`
 - `pshared`: allow the semaphore to be shared by multiple processes
 - General situations: set `pshared` to 0
 - `value`: initialization value for the semaphore
- `int sem_wait(sem_t * sem);`
- `int sem_post(sem_t * sem);`
- Other: `sem_getvalue`, `sem_destroy`, `sem_trywait`

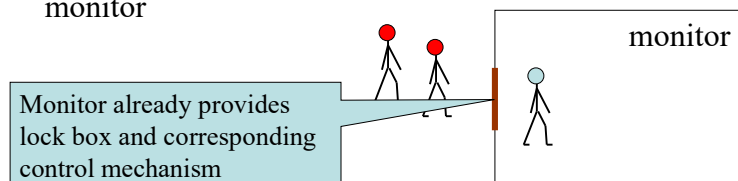
❖ Project 3

Properties of Semaphores

- ❖ Semaphore is a powerful primitive
 - Can do all types of synchronizations
 - however, similar to goto statement → unstructured
 - may issue wait and forget to signal → lock out other processes
 - may reverse the order of multiple waits → deadlock
- ❖ Hard to debug semaphore code
 - Wait and signal of semaphores may spread all over
 - Lack of encapsulation

Monitor

- ❖ Protect shared objects within an abstraction
 - Provide **encapsulation**
 - Accesses to a shared object is confined within a monitor
 - Easier to debug the code
- ❖ Provide mutual exclusive accesses
 - No two processes can be active at the same time within a monitor



Shared Device Program with Monitor

❖ N devices in the system

- Use any of them as long as it is free

monitor mutex_devices:

free: **array** [0..N-1] **of** *boolean*; -- initialized to *true*

int *acquire* () {

for *i* := 0 **to** N-1 **do**

if (*free*[*i*]) **then** { *free*[*i*] := *false*; **return** (*i*); }

return (-1); }

Modify this to always return a free printer

void *release* (*index*: *integer*)

{ *free*[*index*] := *true*; }

Synchronization within a Monitor

❖ Monitor guarantees mutual exclusive accesses

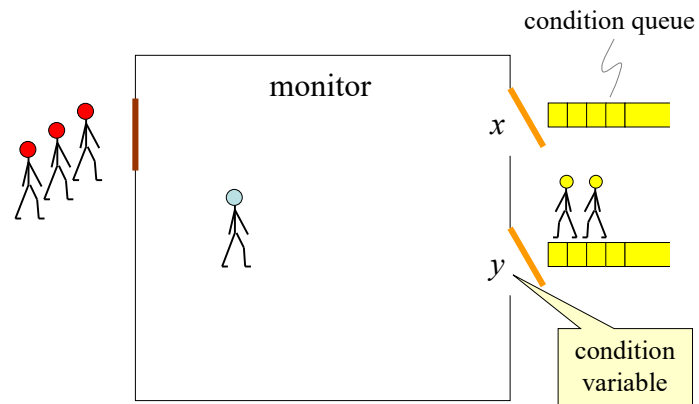
❖ May still need synchronization

- E.g., bounded buffer problem
 - wait for buffer to become full/empty

❖ Monitor also provides condition variables

- To achieve conditional wait
- Associated with each condition variable
 - A condition queue
 - The *wait* and *signal* functions
- Wait inside a monitor
 - Same as wait inside a lockbox protected by a semaphore (mutex)
 - Has potential of deadlock
 - Monitor provides mechanism to counter it

Synchronization within a Monitor



Bounded Buffer Problem with Monitor

monitor bounded_buffer

buffer: **array** [0..*n*-1] **of** *item*;
in, *out*, *counter*: *integer* := 0;
empty, *full*: *condition*;

In semaphore: P (full/empty)
 In monitor: condition wait
 ??? Won't work

function *deposit* (*item*)

```
{ full.wait;
  access buffer;
  counter := counter + 1;
  empty.signal;
}
```

function *remove* (&*item*)

```
{ empty.wait;
  access buffer;
  counter := counter - 1;
  full.signal;
}
```

Bounded Buffer Problem with Monitor

monitor bounded_buffer

buffer: **array** [0..*n*-1] **of** *item*;
in, *out*, *counter*: *integer* := 0;
empty, *full*: *condition*;

In monitor:
 Condition wait/signal are
 memoryless!!!

function *deposit* (*item*)

```
{ if (counter = n) then
    full.wait;
  access buffer;
  counter := counter + 1;
  empty.signal;
}
```

function *remove* (&*item*)

```
{ if (counter = 0) then
    empty.wait;
  access buffer;
  counter := counter - 1;
  full.signal;
}
```

Bounded Buffer Problem with Monitor

Producer process:

```
repeat
  produce item
  deposit (item);
until false;
```

Consumer process:

```
repeat
  remove (item);
  consume item;
until false;
```

Synchronization within a Monitor

❖ Condition variables do not have counting ability

➤ *cond.wait*:

- The calling process always get blocked
- Have to check the condition before doing *cond.wait*

➤ *cond.signal*:

- If no waiting process, the signal is discarded

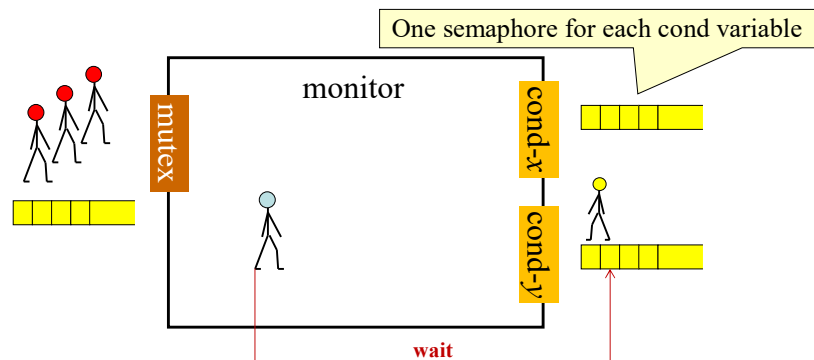
❖ Semaphore has count (which controls sync)

➤ Wait

- Wait may not block the calling process, depending on count
- Can do wait directly without checking

➤ Signal is like being “recorded” since it increases count

Implementing Monitor with Semaphores



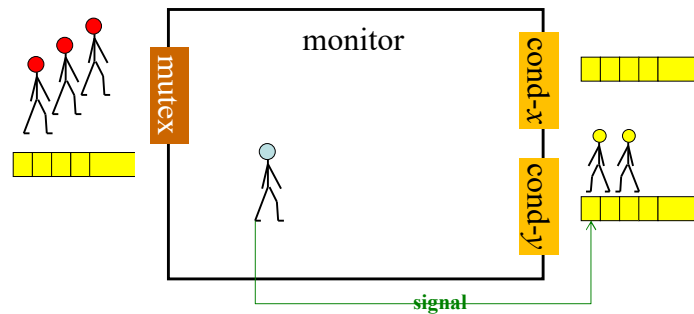
Semaphore “mutex” controls the monitor mutual exclusive entry
 One binary semaphore for each condition variable

Implementation of *cond.wait*: `signal(mutex); wait(cond);`

-- This will be modified later

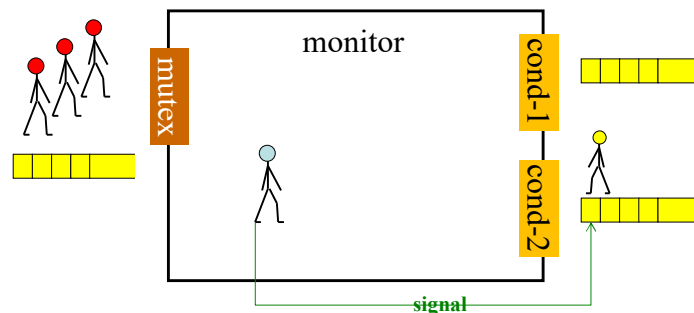
-- Also need to maintain counters to ensure correct implementation

Implementing Monitor with Semaphores



Monitor can only allow one active process within it.
Upon signal, who should be kicked out of the monitor,
signaler or signalee?

Implementing Monitor with Semaphores

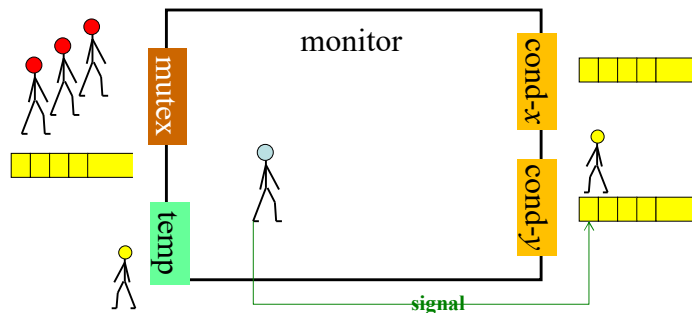


Choose to let the signaler continue after `cond1.signal`
But condition `cond1` may be changed by the signaler's subsequent code
⇒ when signalee comes in monitor, the wait condition should be checked again
(put extra burden on the programmer)

Benefit:

- If signaler code does not change the condition ⇒ Save on context switch
- If signaler is going to exit ⇒ Great, no need to incur the queuing overhead

Implementing Monitor with Semaphores



Choose to let the signalee continue after `cond1.signal` (common choice)

Signaler: moved to temp queue by `temp.wait`

Every process upon exiting the monitor \Rightarrow release signaler first

\Rightarrow `if temp.count > 0, temp.signal; else mutex.signal`

*** temp keeps all waiting signalers (no need to be cond specific)

Monitor and Semaphore

❖ Monitor provides encapsulation

- What should be encapsulated?
- Too much \Rightarrow reduce concurrency
 - Some part of the code that can be executed concurrently, if encapsulated in the monitor, can cause reduced concurrency
 - If not encapsulate them, then lose the meaning of encapsulation
 - Reader/writer example

❖ If monitor is used to do what a semaphore would do

- Monitor is more expensive

Monitor Deficiency

❖ Reduce concurrency

```

monitor rw_object
{ shared_object: some type;

    function read ()
    { read and return some attributes of shared_object; }

    function write (x)
    { write x to shared_object; }
}

```

Monitor Deficiency

```

Monitor rw_object
{ rcount: integer; busy: boolean;
  ok2read, ok2write: condition;

  procedure startread;
  if busy then
    ok2read.wait;
    rcount := rcount + 1;
    ok2read.signal
  end;

  procedure startwrite;
  if busy or rcount != 0 then
    ok2write.wait;
    busy := true;
  end;

  procedure endread;
  rcount := rcount - 1;
  if rcount = 0 then
    ok2write.signal;
  end;

  procedure endwrite;
  busy := false;
  if empty (ok2read.queue) then
    ok2write.signal;
  else ok2read.signal;
  end;
}

```

CS is outside \Rightarrow Does not provide sufficient encapsulation!
 \Rightarrow Defeat the purpose of monitor

Readings

❖ All of Chapter 5